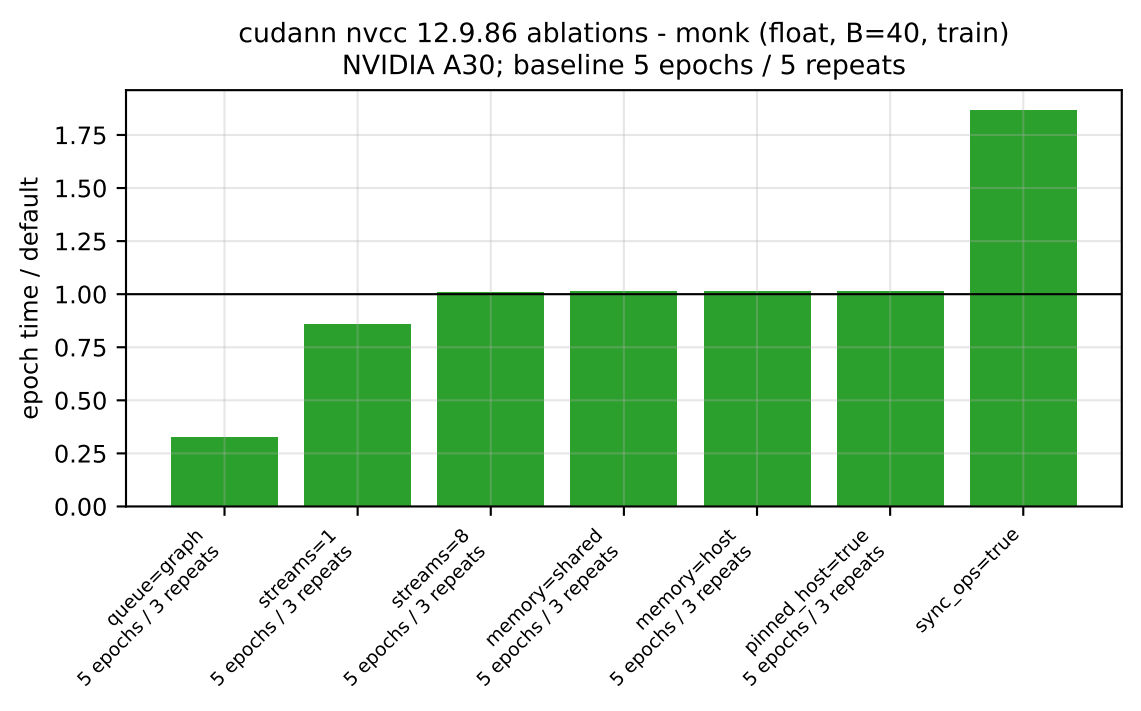
cudann nvcc 12.9.86 ablations - monk (float, B=40, train)
NVIDIA A30; baseline 5 epochs / 5 repeats
epoch time / default
1.75
1.50
1.25
1.00
0.75
0.50
0.25
0.00
queue=graph
5 epochs / 3 repeats
streams=1
5 epochs / 3 repeats
streams=8
5 epochs / 3 repeats
memory=shared
5 epochs / 3 repeats
memory=host
5 epochs / 3 repeats
pinned_host=true
5 epochs / 3 repeats
sync_ops=true