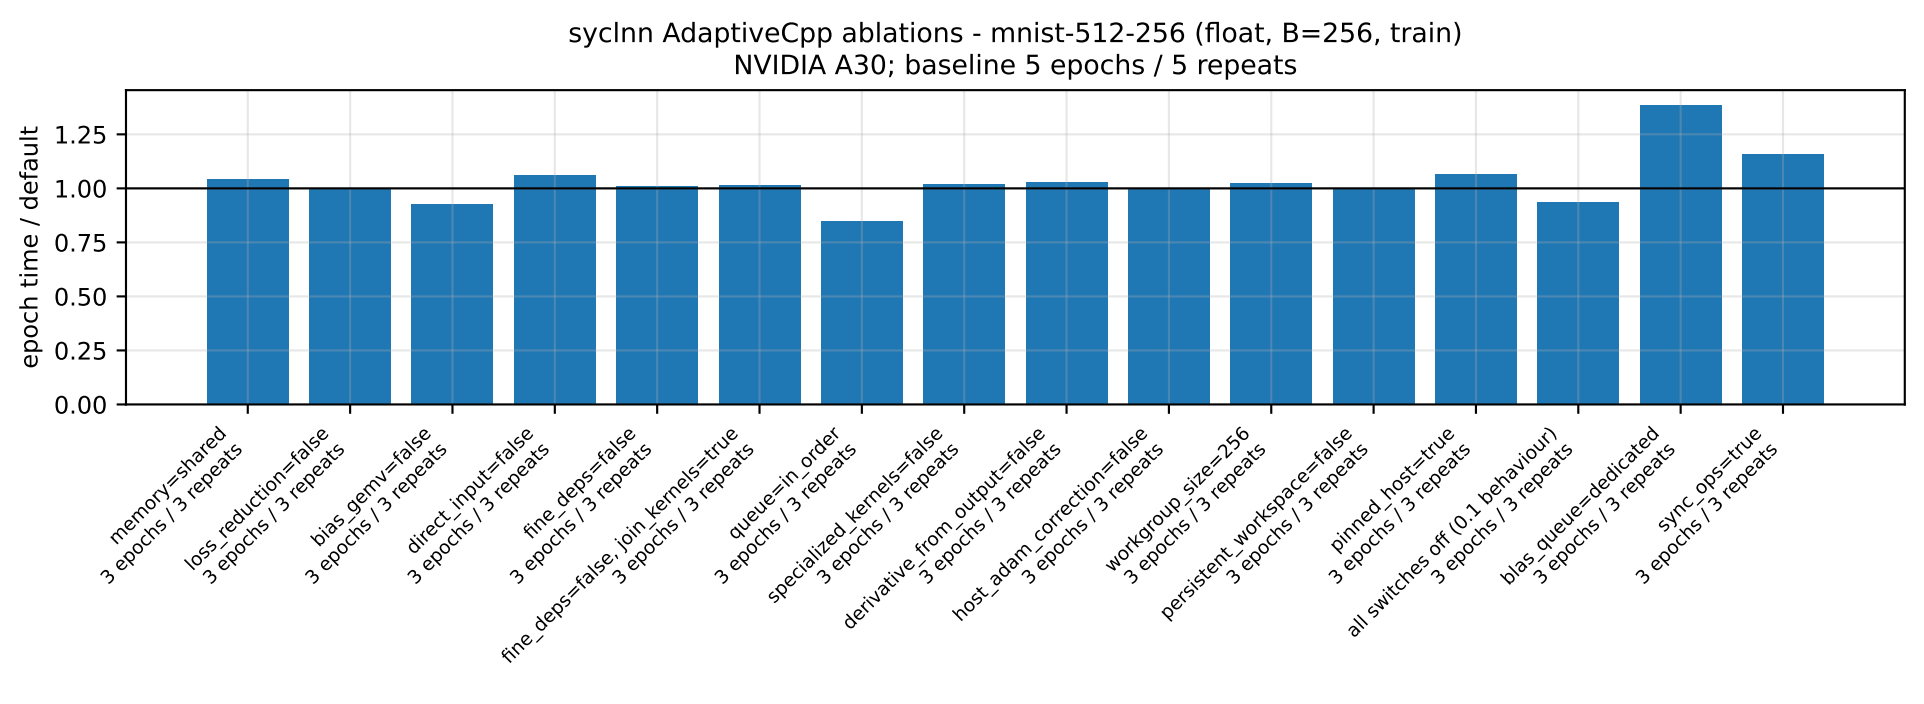
sycInn AdaptiveCpp ablations - mnist-512-256 (float, B=256, train)
NVIDIA A30; baseline 5 epochs / 5 repeats
epoch time / default
0.00
0.25
0.50
0.75
1.00
1.25
memory=shared
3 epochs / 3 repeats
loss_reduction=false
3 epochs / 3 repeats
bias_gemv=false
3 epochs / 3 repeats
direct_input=false
3 epochs / 3 repeats
fine_deps=false
3 epochs / 3 repeats
fine_deps=false, join_kernels=true
3 epochs / 3 repeats
queue=in_order
3 epochs / 3 repeats
specialized_kernels=false
3 epochs / 3 repeats
derivative_from_output=false
3 epochs / 3 repeats
host_adam_correction=false
3 epochs / 3 repeats
workgroup_size=256
3 epochs / 3 repeats
persistent_workspace=false
3 epochs / 3 repeats
pinned_host=true
3 epochs / 3 repeats
all switches off (0.1 behaviour)
3 epochs / 3 repeats
blas_queue=dedicated
3 epochs / 3 repeats
sync_ops=true
3 epochs / 3 repeats