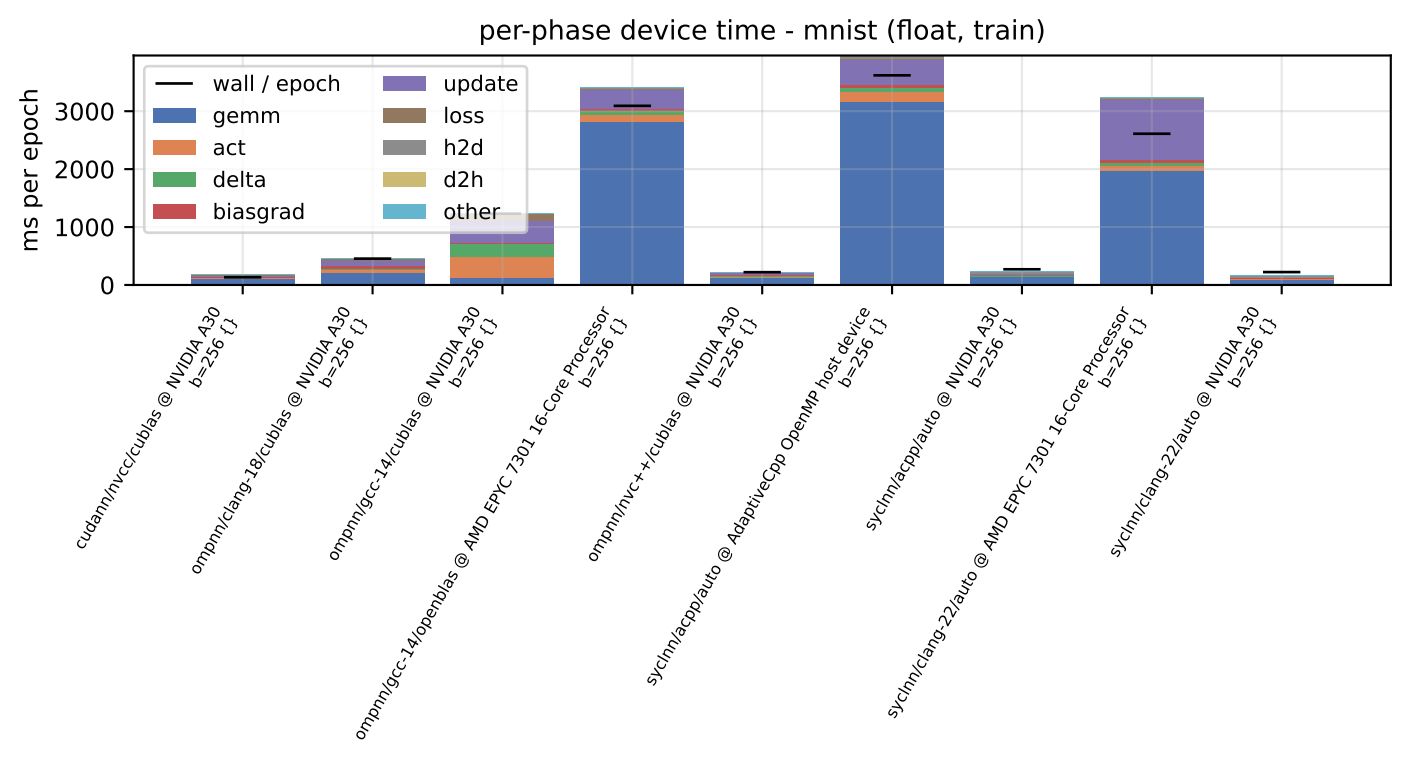
per-phase device time - mnist (float, train)
ms per epoch
0
1000
2000
3000
wall / epoch
gemm
act
delta
biasgrad
update
loss
h2d
d2h
other
cudann/nvcc/cublas @ NVIDIA A30 b=256 {}
ompnn/clang-18/cublas @ NVIDIA A30 b=256 {}
ompnn/gcc-14/cublas @ NVIDIA A30 b=256 {}
ompnn/gcc-14/openblas @ AMD EPYC 7301 16-Core Processor b=256 {}
ompnn/nvc++/cublas @ NVIDIA A30 b=256 {}
syclnn/acpp/auto @ AdaptiveCpp OpenMP host device b=256 {}
syclnn/acpp/auto @ NVIDIA A30 b=256 {}
syclnn/clang-22/auto @ AMD EPYC 7301 16-Core Processor b=256 {}
syclnn/clang-22/auto @ NVIDIA A30 b=256 {}