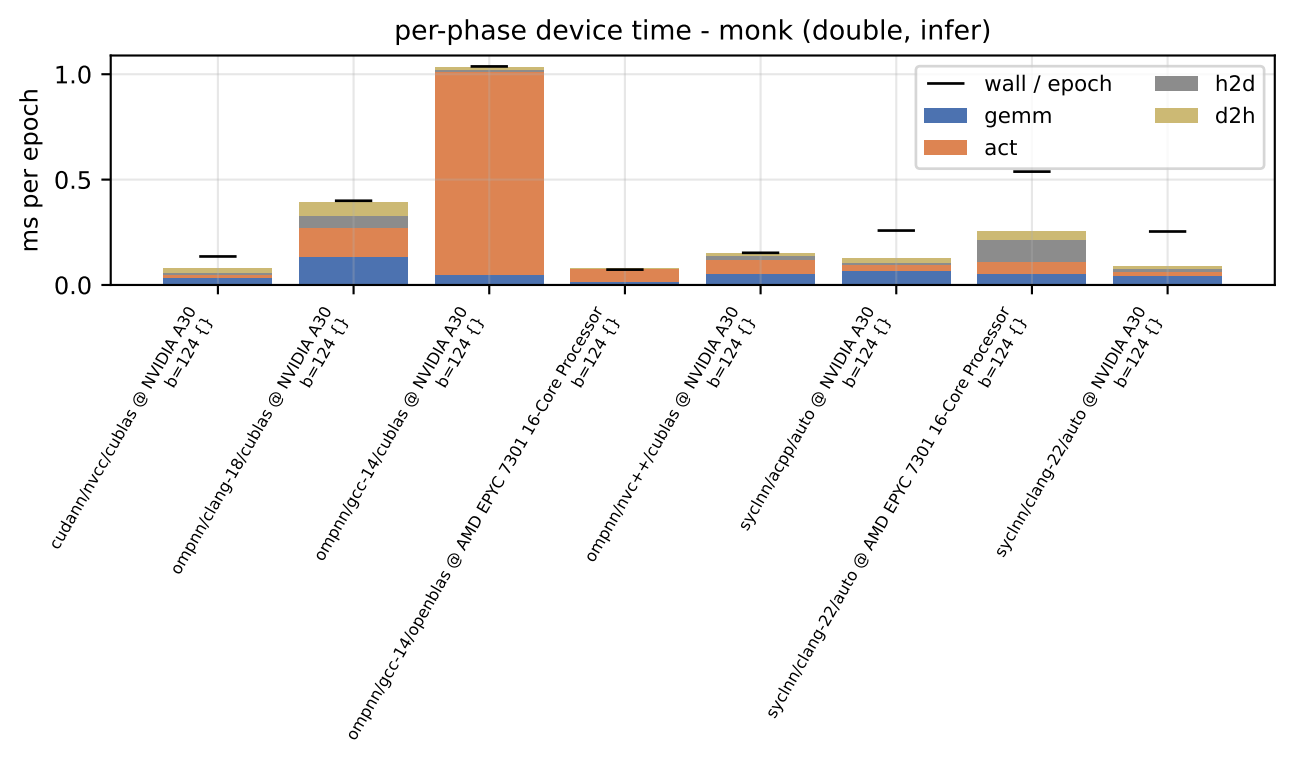
per-phase device time - monk (double, infer)
ms per epoch
1.0
0.5
0.0
wall / epoch
gemm
act
h2d
d2h
cudann/nvcc/cublas @ NVIDIA A30 b=124 {}
ompnn/clang-18/cublas @ NVIDIA A30 b=124 {}
ompnn/gcc-14/cublas @ NVIDIA A30 b=124 {}
ompnn/gcc-14/openblas @ AMD EPYC 7301 16-Core Processor b=124 {}
ompnn/nvc++/cublas @ NVIDIA A30 b=124 {}
syclnn/acpp/auto @ NVIDIA A30 b=124 {}
syclnn/clang-22/auto @ AMD EPYC 7301 16-Core Processor b=124 {}
syclnn/clang-22/auto @ NVIDIA A30 b=124 {}