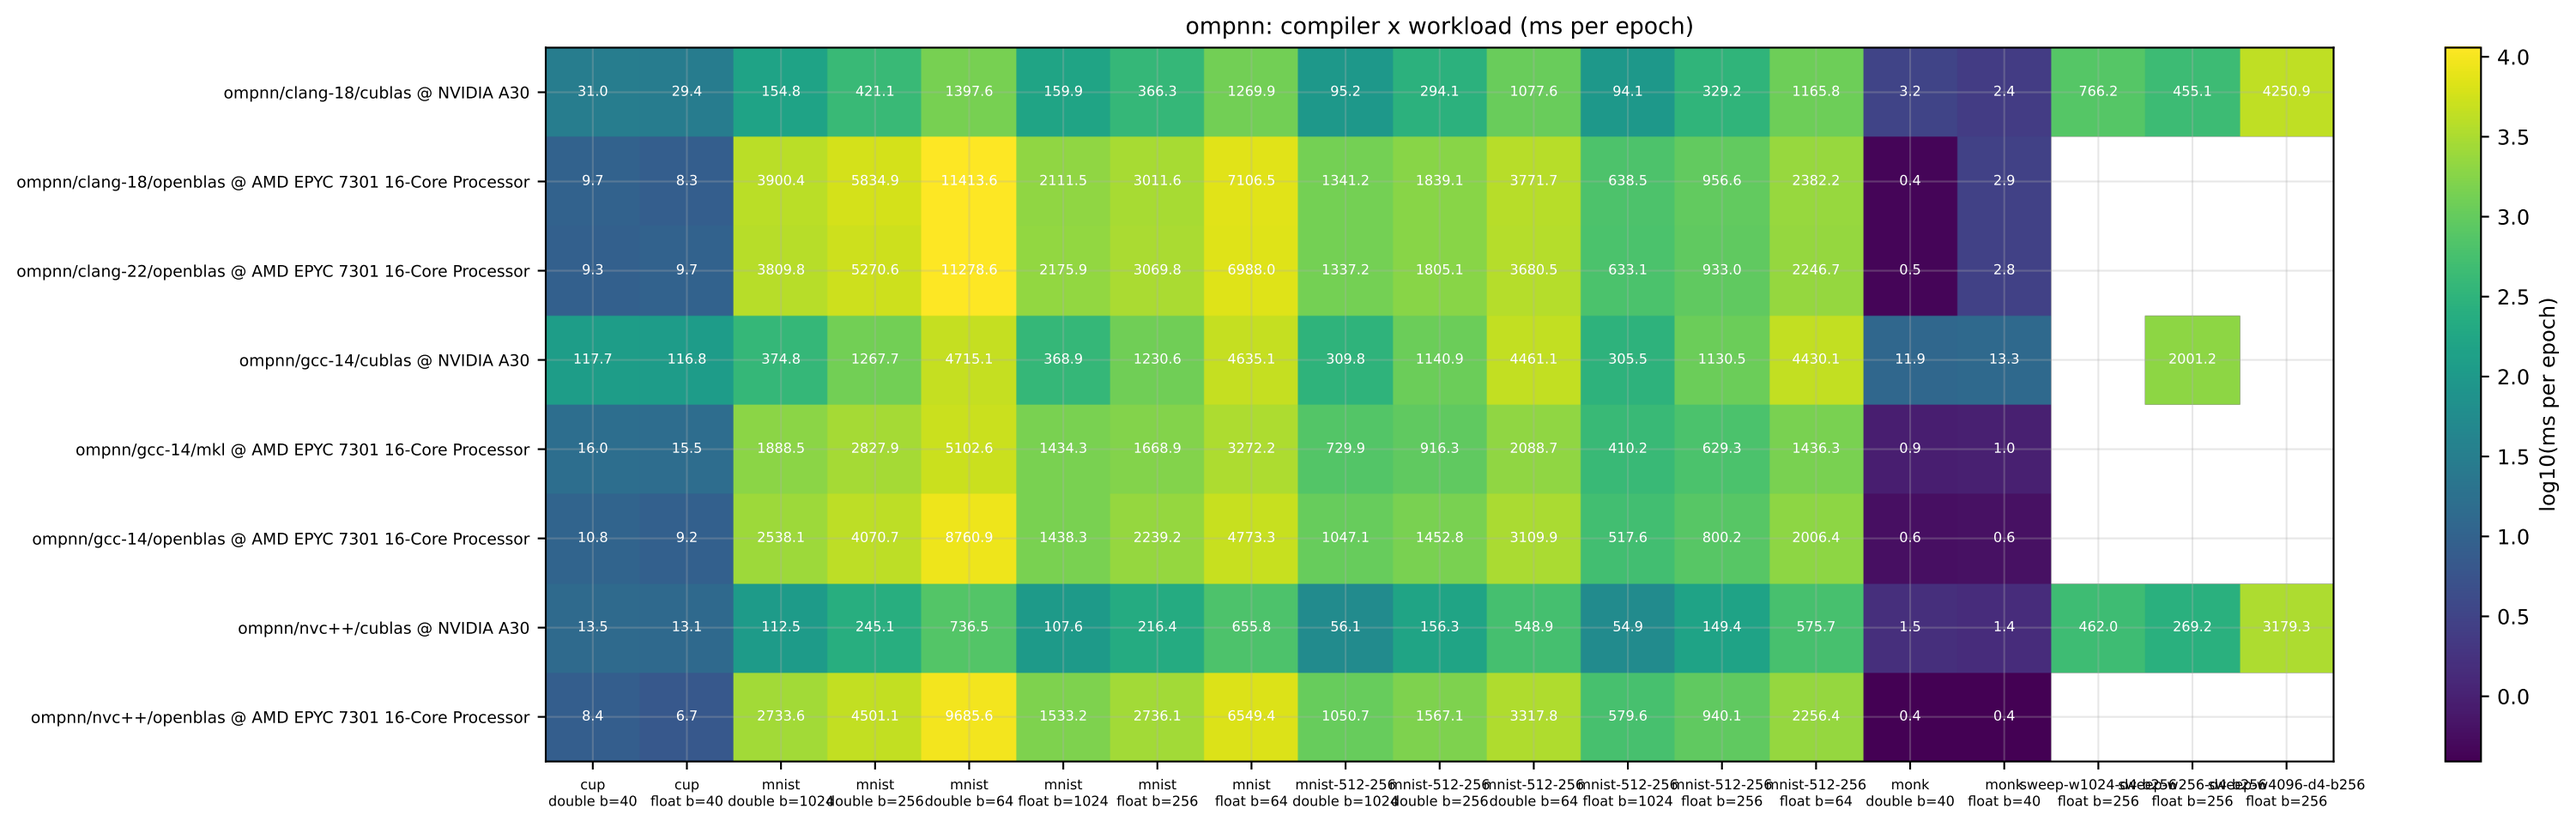
ompnn: compiler x workload (ms per epoch)
ompnn/clang-18/cublas @ NVIDIA A30
ompnn/clang-18/openblas @ AMD EPYC 7301 16-Core Processor
ompnn/clang-22/openblas @ AMD EPYC 7301 16-Core Processor
ompnn/gcc-14/cublas @ NVIDIA A30
ompnn/gcc-14/mkl @ AMD EPYC 7301 16-Core Processor
ompnn/gcc-14/openblas @ AMD EPYC 7301 16-Core Processor
ompnn/nvc++/cublas @ NVIDIA A30
ompnn/nvc++/openblas @ AMD EPYC 7301 16-Core Processor
31.0 29.4 154.8 421.1 1397.6 159.9 366.3 1269.9 95.2 294.1 1077.6 94.1 329.2 1165.8 3.2 2.4 766.2 455.1 4250.9
9.7 8.3 3900.4 5834.9 11413.6 2111.5 3011.6 7106.5 1341.2 1839.1 3771.7 638.5 956.6 2382.2 0.4 2.9
9.3 9.7 3809.8 5270.6 11278.6 2175.9 3069.8 6988.0 1337.2 1805.1 3680.5 633.1 933.0 2246.7 0.5 2.8
117.7 116.8 374.8 1267.7 4715.1 368.9 1230.6 4635.1 309.8 1140.9 4461.1 305.5 1130.5 4430.1 11.9 13.3 2001.2
16.0 15.5 1888.5 2827.9 5102.6 1434.3 1668.9 3272.2 729.9 916.3 2088.7 410.2 629.3 1436.3 0.9 1.0
10.8 9.2 2538.1 4070.7 8760.9 1438.3 2239.2 4773.3 1047.1 1452.8 3109.9 517.6 800.2 2006.4 0.6 0.6
13.5 13.1 112.5 245.1 736.5 107.6 216.4 655.8 56.1 156.3 548.9 54.9 149.4 575.7 1.5 1.4 462.0 269.2 3179.3
8.4 6.7 2733.6 4501.1 9685.6 1533.2 2736.1 6549.4 1050.7 1567.1 3317.8 579.6 940.1 2256.4 0.4 0.4
cup double b=40
cup float b=40
mnist double b=1024
mnist double b=256
mnist double b=64
mnist float b=1024
mnist float b=256
mnist float b=64
mnist-512-256 double b=1024
mnist-512-256 double b=256
mnist-512-256 double b=64
mnist-512-256 float b=1024
mnist-512-256 float b=256
mnist-512-256 float b=64
monk double b=40
monk float b=40
sweep-w1024-d4-b256
sweep-w256-d4-b256
sweep-w4096-d4-b256
float b=256
float b=256
float b=256
log10(ms per epoch)
4.0
3.5
3.0
2.5
2.0
1.5
1.0
0.5
0.0