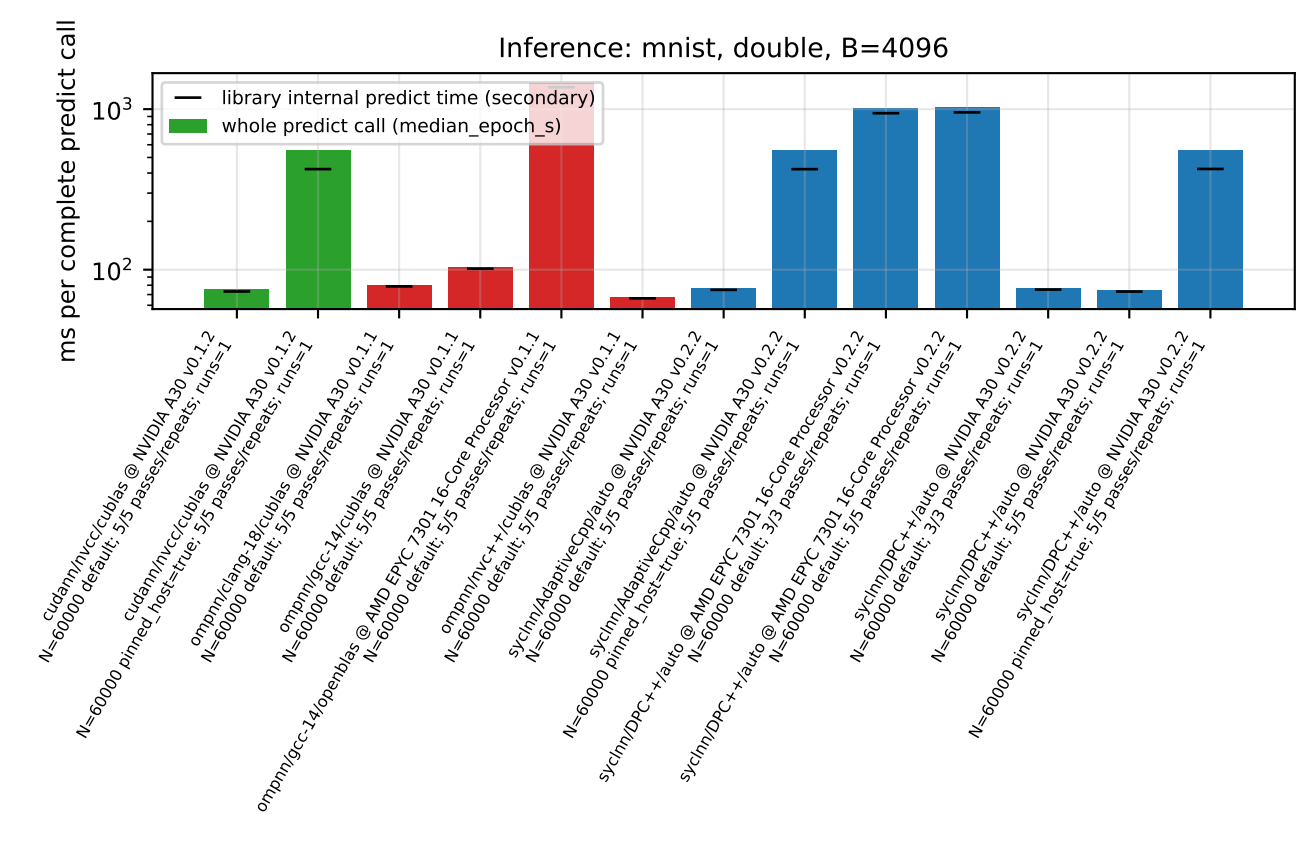
Inference: mnist, double, B=4096
ms per complete predict call
$10^3$
$10^2$
library internal predict time (secondary)
whole predict call (median_epoch_s)
cudann/nvcc/cublas @ NVIDIA A30 v0.1.2 N=60000 default; 5/5 passes/repeats; runs=1
cudann/nvcc/cublas @ NVIDIA A30 v0.1.2 N=60000 pinned_host=true; 5/5 passes/repeats; runs=1
ompnn/clang-18/cublas @ NVIDIA A30 v0.1.1 N=60000 default; 5/5 passes/repeats; runs=1
ompnn/gcc-14/cublas @ NVIDIA A30 v0.1.1 N=60000 default; 5/5 passes/repeats; runs=1
ompnn/gcc-14/openblas @ AMD EPYC 7301 16-Core Processor v0.1.1 N=60000 default; 5/5 passes/repeats; runs=1
ompnn/nvc++/cublas @ NVIDIA A30 v0.1.1 N=60000 default; 5/5 passes/repeats; runs=1
syclnn/AdaptiveCpp/auto @ NVIDIA A30 v0.2.2 N=60000 default; 5/5 passes/repeats; runs=1
syclnn/AdaptiveCpp/auto @ NVIDIA A30 v0.2.2 N=60000 pinned_host=true; 5/5 passes/repeats; runs=1
syclnn/DPC++/auto @ AMD EPYC 7301 16-Core Processor v0.2.2 N=60000 default; 3/3 passes/repeats; runs=1
syclnn/DPC++/auto @ AMD EPYC 7301 16-Core Processor v0.2.2 N=60000 default; 5/5 passes/repeats; runs=1
syclnn/DPC++/auto @ NVIDIA A30 v0.2.2 N=60000 default; 3/3 passes/repeats; runs=1
syclnn/DPC++/auto @ NVIDIA A30 v0.2.2 N=60000 default; 5/5 passes/repeats; runs=1
syclnn/DPC++/auto @ NVIDIA A30 v0.2.2 N=60000 pinned_host=true; 5/5 passes/repeats; runs=1