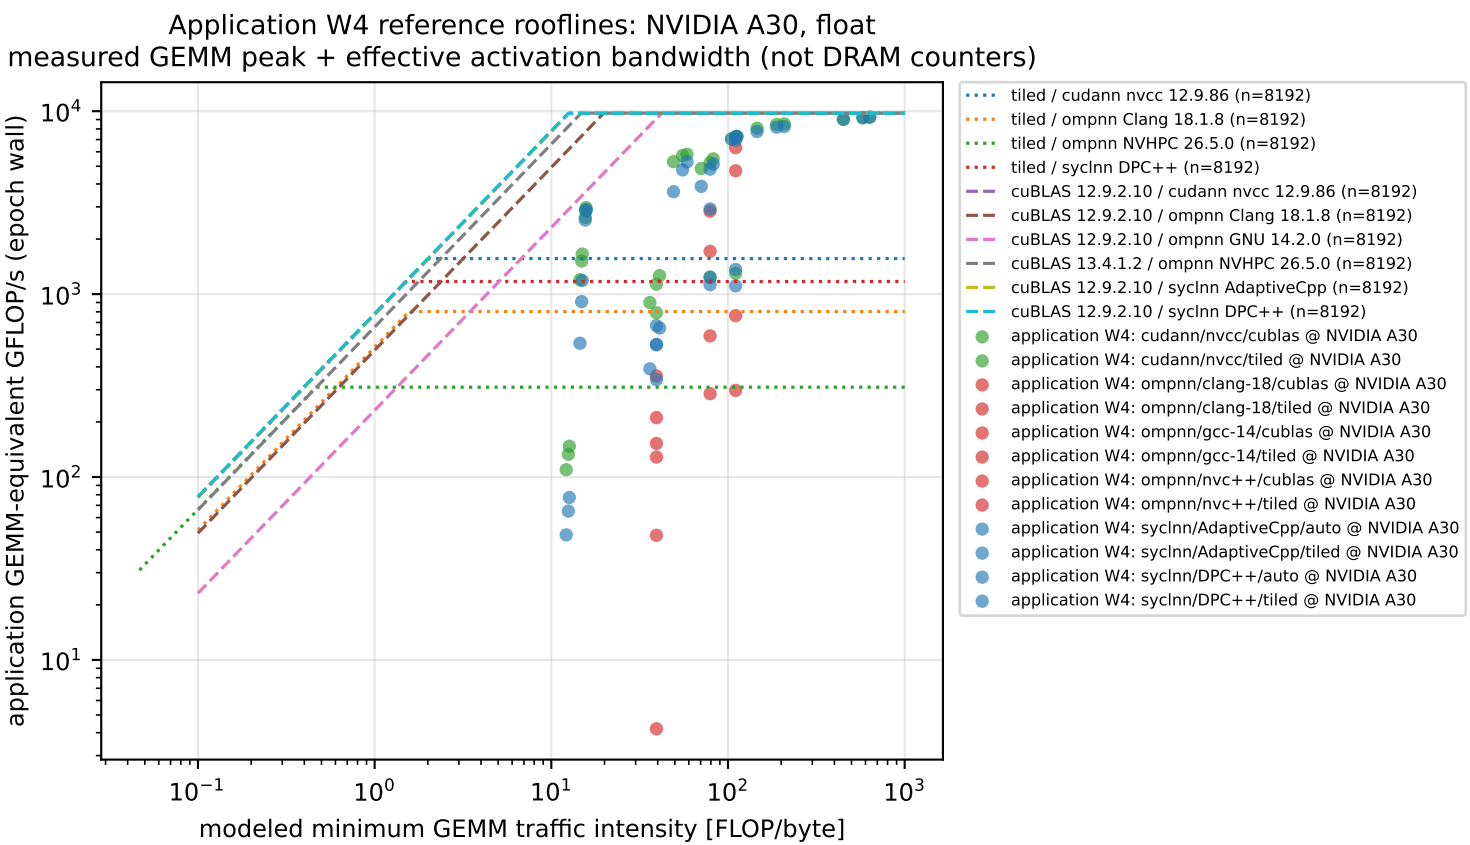
Application W4 reference rooflines: NVIDIA A30, float
measured GEMM peak + effective activation bandwidth (not DRAM counters)
application GEMM-equivalent GFLOP/s (epoch wall)
modeled minimum GEMM traffic intensity [FLOP/byte]
tiled / cudann nvcc 12.9.86 (n=8192)
tiled / ompnn Clang 18.1.8 (n=8192)
tiled / ompnn NVHPC 26.5.0 (n=8192)
tiled / syclnn DPC++ (n=8192)
cuBLAS 12.9.2.10 / cudann nvcc 12.9.86 (n=8192)
cuBLAS 12.9.2.10 / ompnn Clang 18.1.8 (n=8192)
cuBLAS 12.9.2.10 / ompnn GNU 14.2.0 (n=8192)
cuBLAS 13.4.1.2 / ompnn NVHPC 26.5.0 (n=8192)
cuBLAS 12.9.2.10 / syclnn AdaptiveCpp (n=8192)
cuBLAS 12.9.2.10 / syclnn DPC++ (n=8192)
application W4: cudann/nvcc/cublas @ NVIDIA A30
application W4: cudann/nvcc/tiled @ NVIDIA A30
application W4: ompnn/clang-18/cublas @ NVIDIA A30
application W4: ompnn/clang-18/tiled @ NVIDIA A30
application W4: ompnn/gcc-14/cublas @ NVIDIA A30
application W4: ompnn/gcc-14/tiled @ NVIDIA A30
application W4: ompnn/nvc++/cublas @ NVIDIA A30
application W4: ompnn/nvc++/tiled @ NVIDIA A30
application W4: syclnn/AdaptiveCpp/auto @ NVIDIA A30
application W4: syclnn/AdaptiveCpp/tiled @ NVIDIA A30
application W4: syclnn/DPC++/auto @ NVIDIA A30
application W4: syclnn/DPC++/tiled @ NVIDIA A30