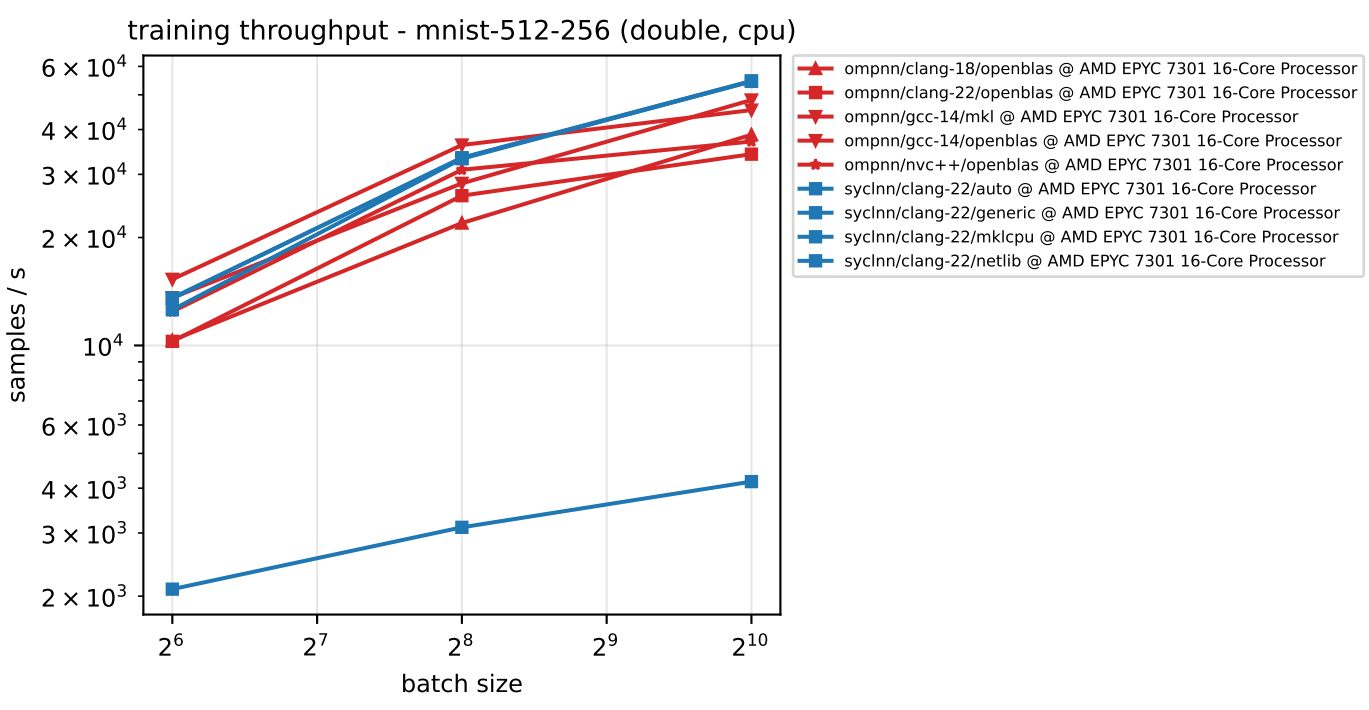
training throughput - mnist-512-256 (double, cpu)
samples / s
batch size
6×10^4
4×10^4
3×10^4
2×10^4
10^4
6×10^3
4×10^3
3×10^3
2×10^3
2^6
2^7
2^8
2^9
2^{10}
ompnn/clang-18/openblas @ AMD EPYC 7301 16-Core Processor
ompnn/clang-22/openblas @ AMD EPYC 7301 16-Core Processor
ompnn/gcc-14/mkl @ AMD EPYC 7301 16-Core Processor
ompnn/gcc-14/openblas @ AMD EPYC 7301 16-Core Processor
ompnn/nvc++/openblas @ AMD EPYC 7301 16-Core Processor
syclnn/clang-22/auto @ AMD EPYC 7301 16-Core Processor
syclnn/clang-22/generic @ AMD EPYC 7301 16-Core Processor
syclnn/clang-22/mklcpu @ AMD EPYC 7301 16-Core Processor
syclnn/clang-22/netlib @ AMD EPYC 7301 16-Core Processor