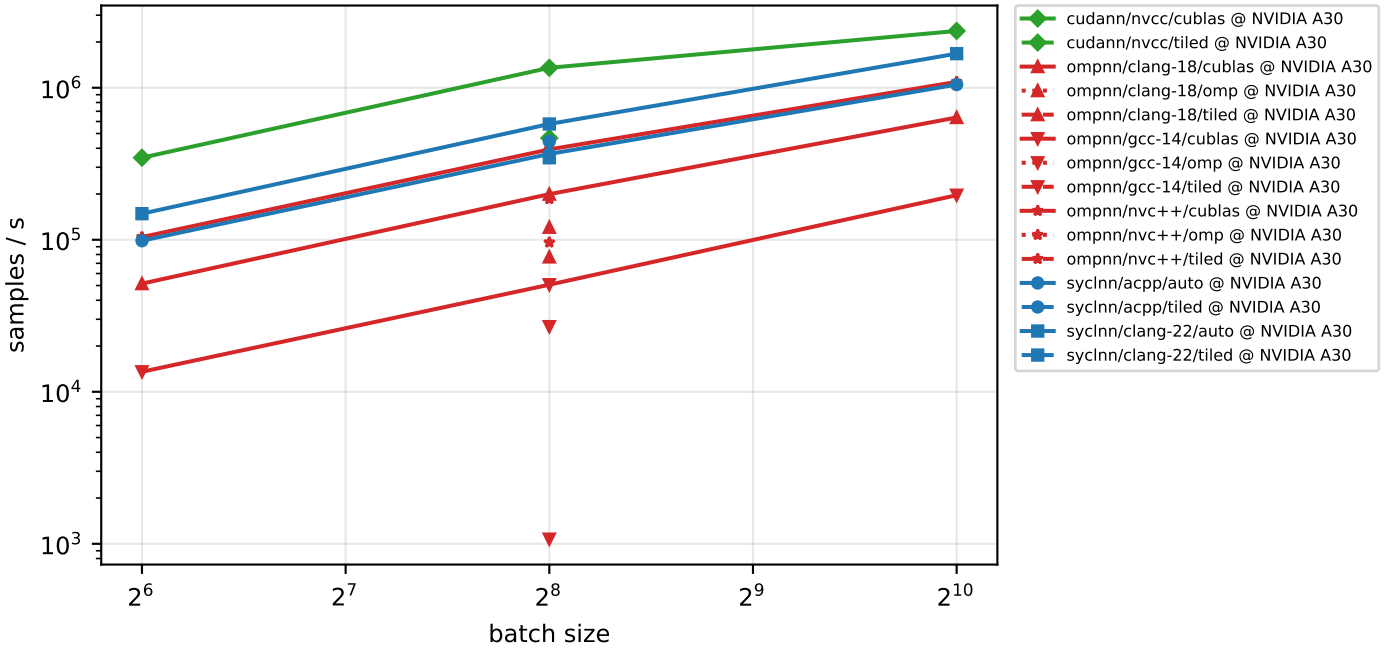
training throughput - mnist-512-256 (float, gpu)
samples / s
batch size
10^6
10^5
10^4
10^3
2^6
2^7
2^8
2^9
2^{10}
cudann/nvcc/cublas @ NVIDIA A30
cudann/nvcc/tiled @ NVIDIA A30
ompnn/clang-18/cublas @ NVIDIA A30
ompnn/clang-18/omp @ NVIDIA A30
ompnn/clang-18/tiled @ NVIDIA A30
ompnn/gcc-14/cublas @ NVIDIA A30
ompnn/gcc-14/omp @ NVIDIA A30
ompnn/gcc-14/tiled @ NVIDIA A30
ompnn/nvc++/cublas @ NVIDIA A30
ompnn/nvc++/omp @ NVIDIA A30
ompnn/nvc++/tiled @ NVIDIA A30
syclnn/acpp/auto @ NVIDIA A30
syclnn/acpp/tiled @ NVIDIA A30
syclnn/clang-22/auto @ NVIDIA A30
syclnn/clang-22/tiled @ NVIDIA A30