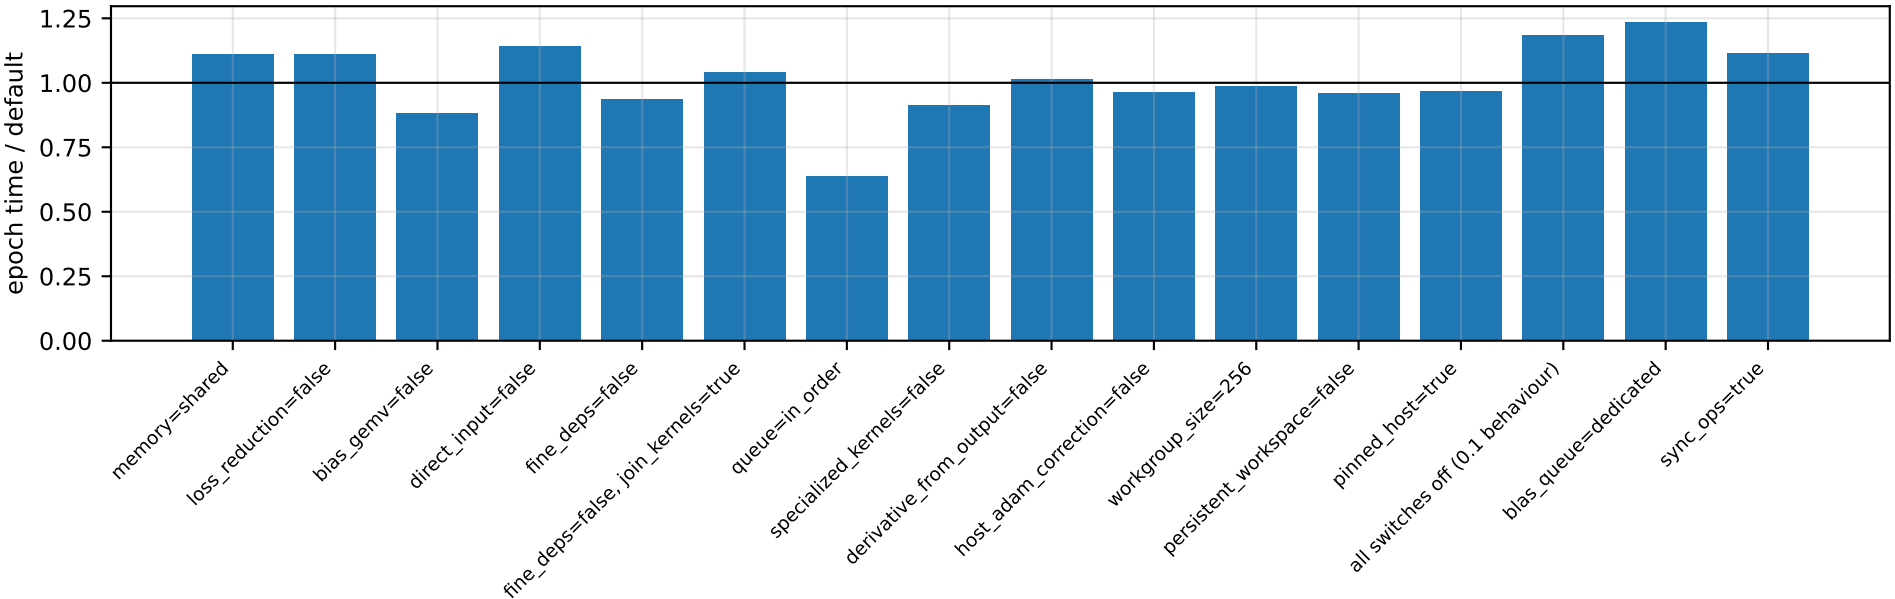
sycInn AdaptiveCpp ablations - cup (float, B=40, train)
AMD Radeon RX 7900 XTX; baseline 5 epochs / 3 repeats
epoch time / default
1.25
1.00
0.75
0.50
0.25
0.00
memory=shared
loss_reduction=false
bias_gemv=false
direct_input=false
fine_deps=false
fine_deps=false, join_kernels=true
queue=in_order
specialized_kernels=false
derivative_from_output=false
host_adam_correction=false
workgroup_size=256
persistent_workspace=false
pinned_host=true
all switches off (0.1 behaviour)
blas_queue=dedicated
sync_ops=true