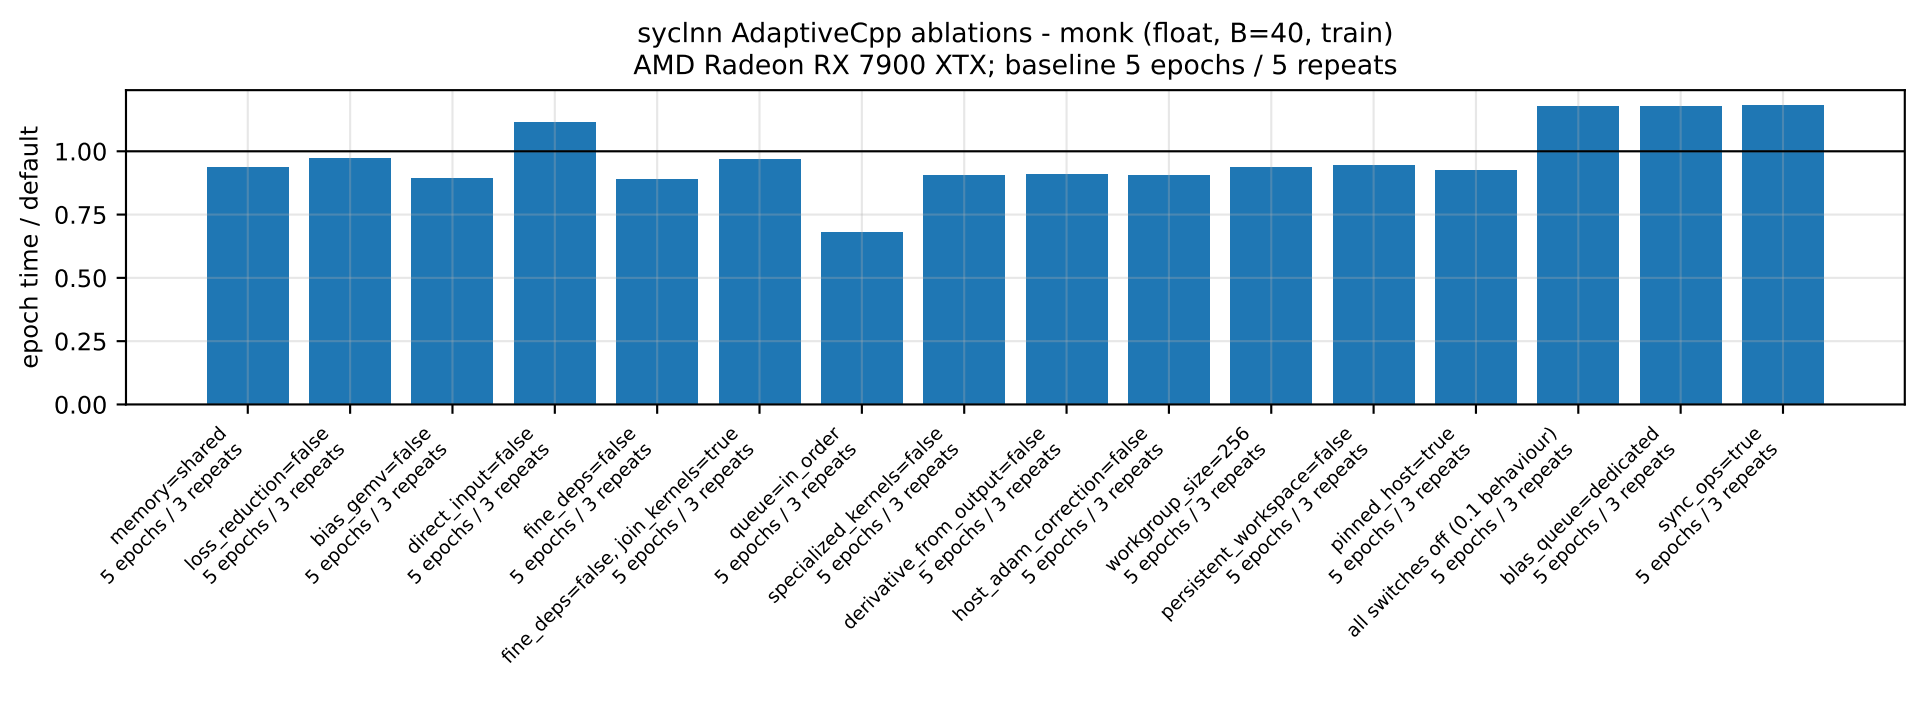
sycInn AdaptiveCpp ablations - monk (float, B=40, train)
AMD Radeon RX 7900 XTX; baseline 5 epochs / 5 repeats
epoch time / default
0.00
0.25
0.50
0.75
1.00
memory=shared
5 epochs / 3 repeats
loss_reduction=false
5 epochs / 3 repeats
bias_gemv=false
5 epochs / 3 repeats
direct_input=false
5 epochs / 3 repeats
fine_deps=false
5 epochs / 3 repeats
fine_deps=false, join_kernels=true
5 epochs / 3 repeats
queue=in_order
5 epochs / 3 repeats
specialized_kernels=false
5 epochs / 3 repeats
derivative_from_output=false
5 epochs / 3 repeats
host_adam_correction=false
5 epochs / 3 repeats
workgroup_size=256
5 epochs / 3 repeats
persistent_workspace=false
5 epochs / 3 repeats
pinned_host=true
5 epochs / 3 repeats
all switches off (0.1 behaviour)
5 epochs / 3 repeats
blas_queue=dedicated
5 epochs / 3 repeats
sync_ops=true
5 epochs / 3 repeats