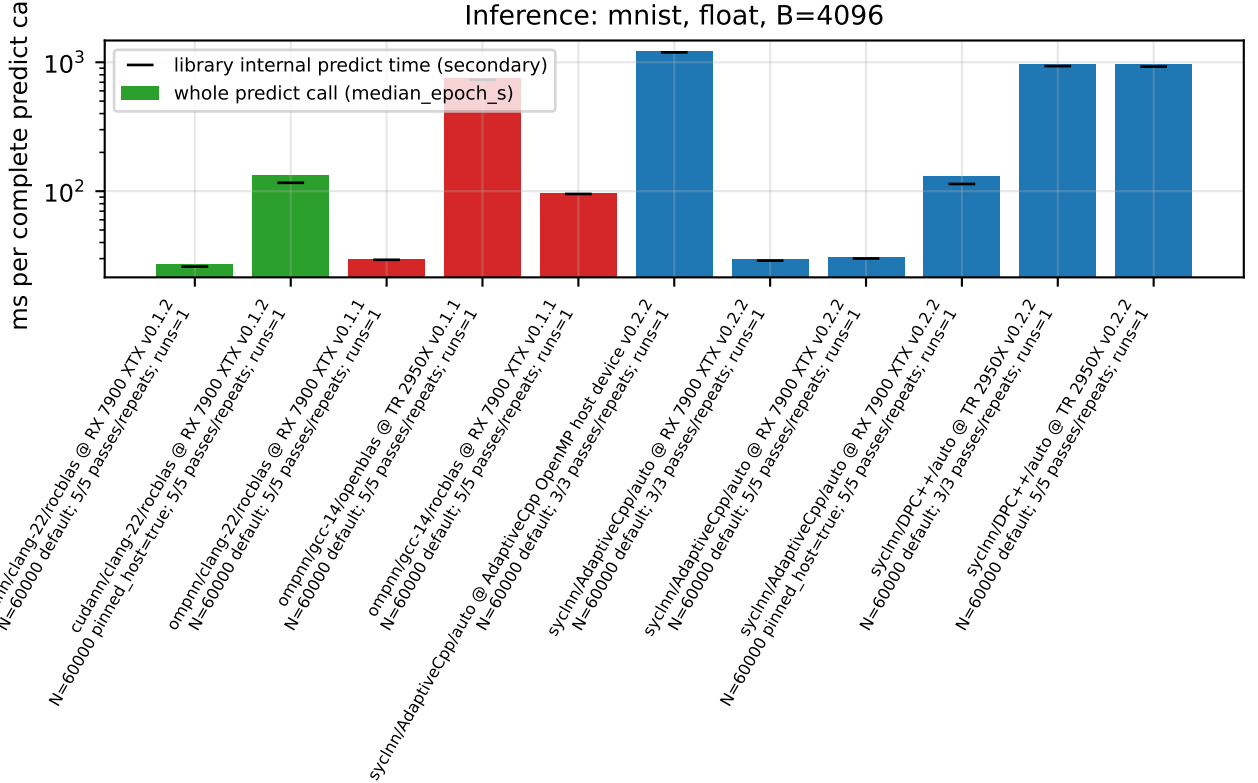
Inference: mnist, float, B=4096
ms per complete predict call
10^3
10^2
library internal predict time (secondary)
whole predict call (median_epoch_s)
cudann/clang-22/rocblas @ RX 7900 XTX v0.1.2 N=60000 default; 5/5 passes/repeats; runs=1
cudann/clang-22/rocblas @ RX 7900 XTX v0.1.2 N=60000 pinned_host=true; 5/5 passes/repeats; runs=1
ompnn/clang-22/rocblas @ RX 7900 XTX v0.1.1 N=60000 default; 5/5 passes/repeats; runs=1
ompnn/gcc-14/openblas @ TR 2950X v0.1.1 N=60000 default; 5/5 passes/repeats; runs=1
ompnn/gcc-14/rocblas @ RX 7900 XTX v0.1.1 N=60000 default; 5/5 passes/repeats; runs=1
syclnn/AdaptiveCpp/auto @ AdaptiveCpp OpenMP host device v0.2.2 N=60000 default; 3/3 passes/repeats; runs=1
syclnn/AdaptiveCpp/auto @ RX 7900 XTX v0.2.2 N=60000 default; 3/3 passes/repeats; runs=1
syclnn/AdaptiveCpp/auto @ RX 7900 XTX v0.2.2 N=60000 default; 5/5 passes/repeats; runs=1
syclnn/AdaptiveCpp/auto @ RX 7900 XTX v0.2.2 N=60000 pinned_host=true; 5/5 passes/repeats; runs=1
syclnn/DPC++/auto @ TR 2950X v0.2.2 N=60000 default; 3/3 passes/repeats; runs=1
syclnn/DPC++/auto @ TR 2950X v0.2.2 N=60000 default; 5/5 passes/repeats; runs=1