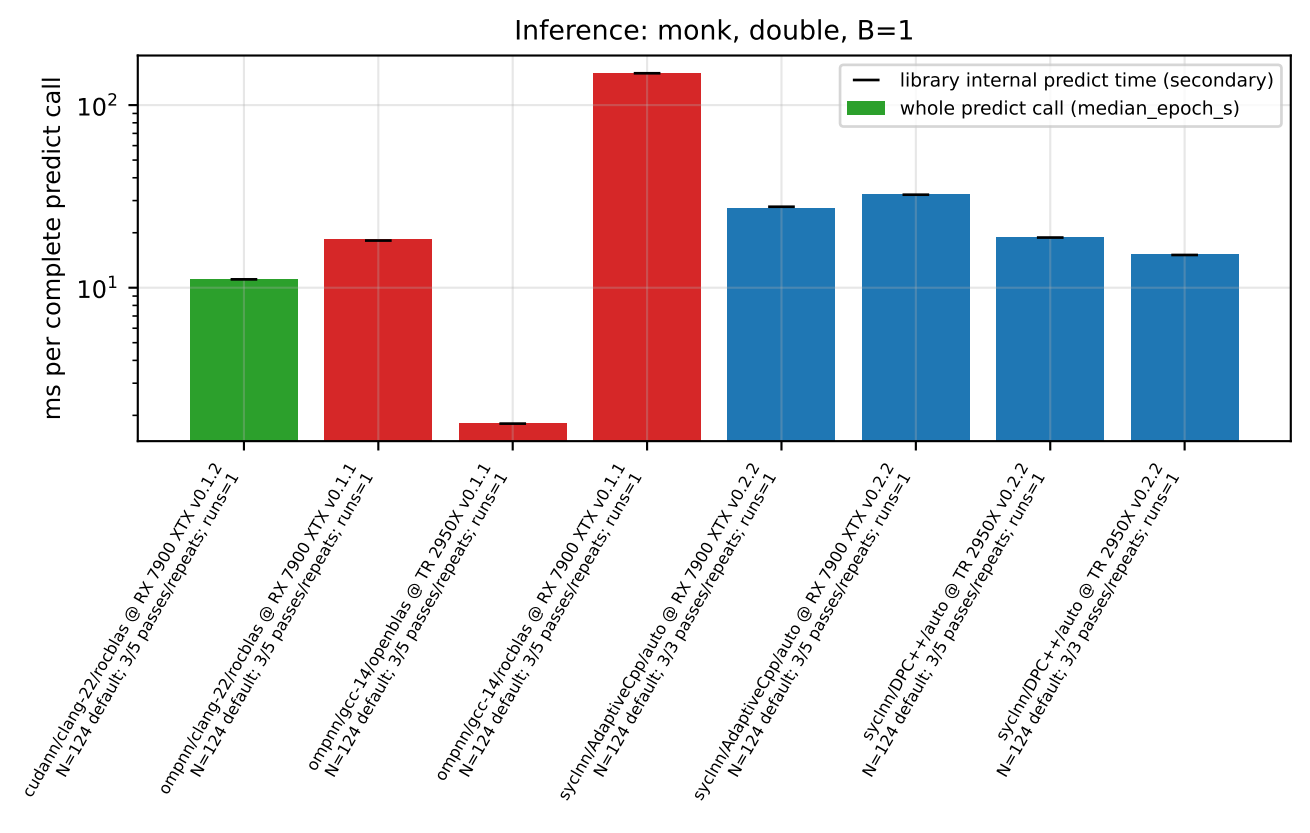
Inference: monk, double, B=1
ms per complete predict call
10^2
10^1
library internal predict time (secondary)
whole predict call (median_epoch_s)
cudann/clang-22/rocblas @ RX 7900 XTX v0.1.2
N=124 default; 3/5 passes/repeats; runs=1
ompnn/clang-22/rocblas @ RX 7900 XTX v0.1.1
N=124 default; 3/5 passes/repeats; runs=1
ompnn/gcc-14/openblas @ TR 2950X v0.1.1
N=124 default; 3/5 passes/repeats; runs=1
ompnn/gcc-14/rocblas @ RX 7900 XTX v0.1.1
N=124 default; 3/5 passes/repeats; runs=1
syclnn/AdaptiveCpp/auto @ RX 7900 XTX v0.2.2
N=124 default; 3/3 passes/repeats; runs=1
syclnn/AdaptiveCpp/auto @ RX 7900 XTX v0.2.2
N=124 default; 3/5 passes/repeats; runs=1
syclnn/DPC++/auto @ TR 2950X v0.2.2
N=124 default; 3/5 passes/repeats; runs=1
syclnn/DPC++/auto @ TR 2950X v0.2.2
N=124 default; 3/3 passes/repeats; runs=1