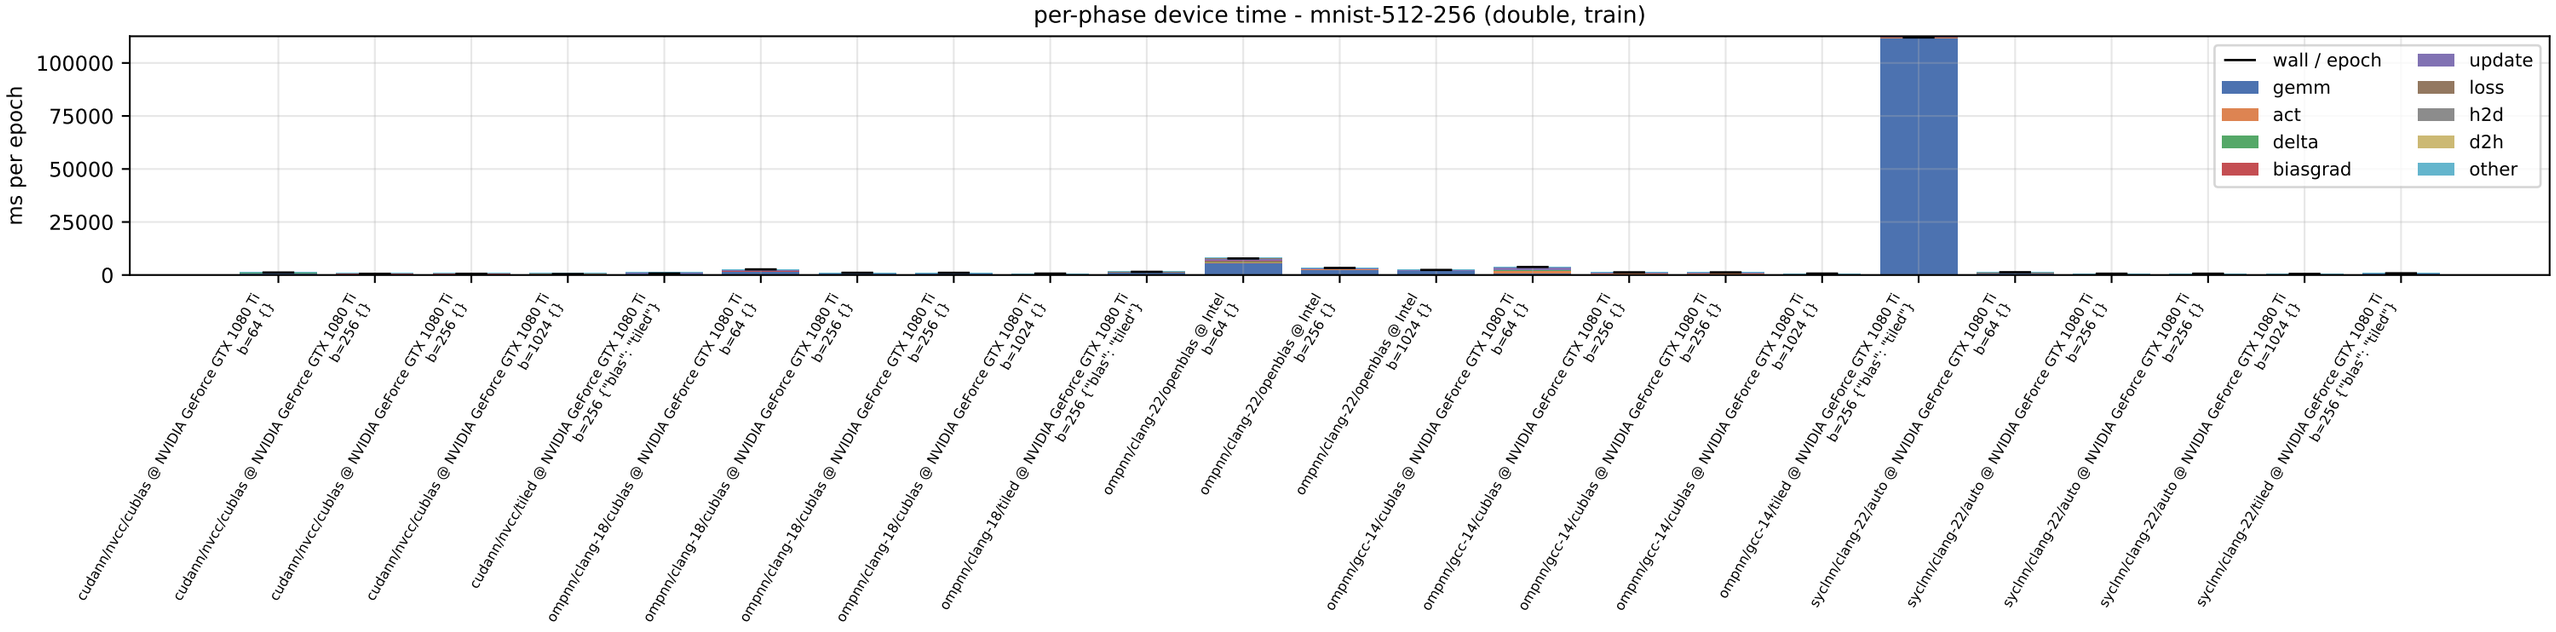

per-phase device time - mnist-512-256 (double, train)
ms per epoch
0
25000
50000
75000
100000
wall / epoch
gemm
act
delta
biasgrad
update
loss
h2d
d2h
other
cudann/nvcc/cublas @ NVIDIA GeForce GTX 1080 Ti b=64 {}
cudann/nvcc/cublas @ NVIDIA GeForce GTX 1080 Ti b=256 {}
cudann/nvcc/cublas @ NVIDIA GeForce GTX 1080 Ti b=256 {}
cudann/nvcc/cublas @ NVIDIA GeForce GTX 1080 Ti b=1024 {}
cudann/nvcc/tiled @ NVIDIA GeForce GTX 1080 Ti b=256 {"bias": "tiled"}
ompnn/clang-18/cublas @ NVIDIA GeForce GTX 1080 Ti b=64 {}
ompnn/clang-18/cublas @ NVIDIA GeForce GTX 1080 Ti b=256 {}
ompnn/clang-18/cublas @ NVIDIA GeForce GTX 1080 Ti b=256 {}
ompnn/clang-18/cublas @ NVIDIA GeForce GTX 1080 Ti b=1024 {}
ompnn/clang-18/tiled @ NVIDIA GeForce GTX 1080 Ti b=256 {"bias": "tiled"}
ompnn/clang-22/openblas @ Intel b=64 {}
ompnn/clang-22/openblas @ Intel b=256 {}
ompnn/clang-22/openblas @ Intel b=1024 {}
ompnn/gcc-14/cublas @ NVIDIA GeForce GTX 1080 Ti b=64 {}
ompnn/gcc-14/cublas @ NVIDIA GeForce GTX 1080 Ti b=256 {}
ompnn/gcc-14/cublas @ NVIDIA GeForce GTX 1080 Ti b=256 {}
ompnn/gcc-14/cublas @ NVIDIA GeForce GTX 1080 Ti b=1024 {}
ompnn/gcc-14/tiled @ NVIDIA GeForce GTX 1080 Ti b=256 {"bias": "tiled"}
syclnn/clang-22/auto @ NVIDIA GeForce GTX 1080 Ti b=64 {}
syclnn/clang-22/auto @ NVIDIA GeForce GTX 1080 Ti b=256 {}
syclnn/clang-22/auto @ NVIDIA GeForce GTX 1080 Ti b=256 {}
syclnn/clang-22/auto @ NVIDIA GeForce GTX 1080 Ti b=1024 {}
syclnn/clang-22/tiled @ NVIDIA GeForce GTX 1080 Ti b=256 {"bias": "tiled"}