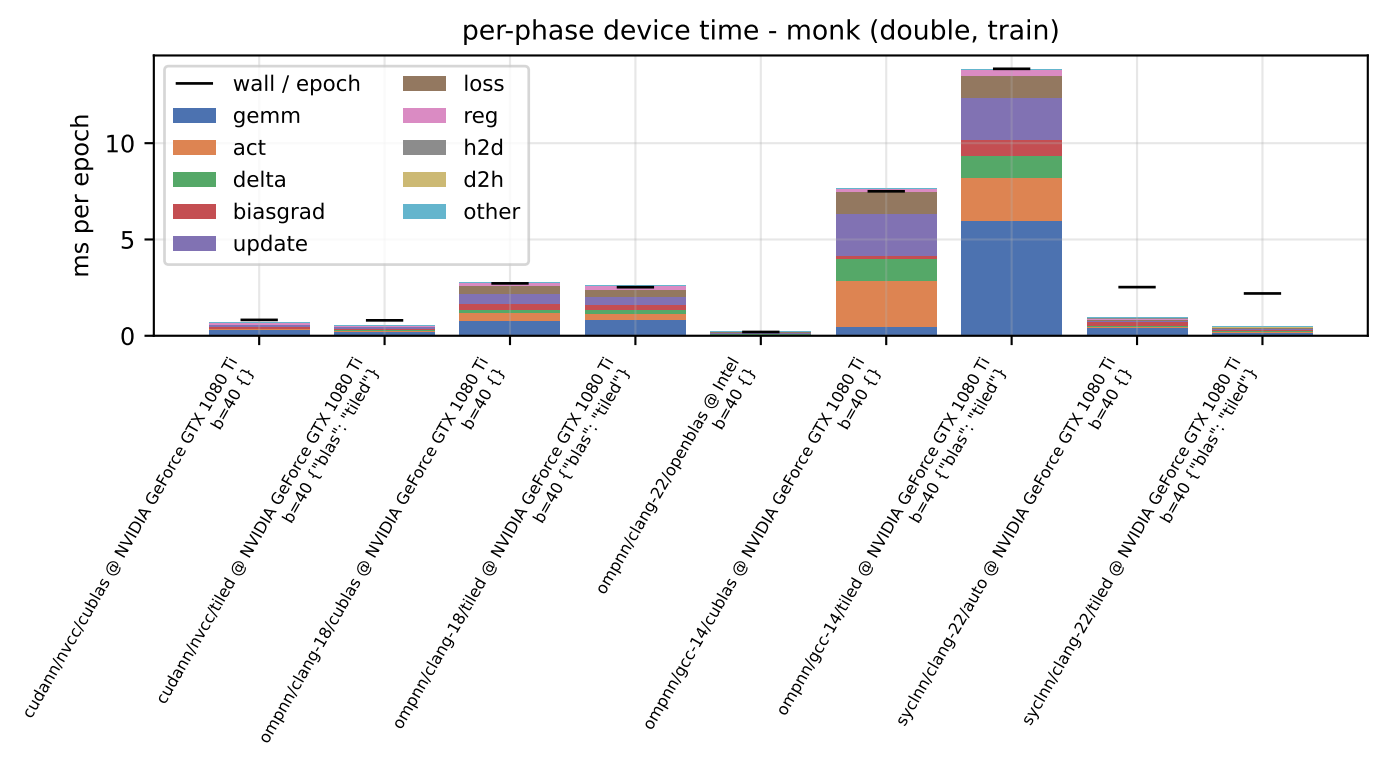

per-phase device time - monk (double, train)
ms per epoch
0
5
10
wall / epoch
gemm
act
delta
biasgrad
update
loss
reg
h2d
d2h
other
cudann/nvcc/cublas @ NVIDIA GeForce GTX 1080 Ti b=40 {}
cudann/nvcc/tiled @ NVIDIA GeForce GTX 1080 Ti b=40 {"blas": "tiled"}
ompnn/clang-18/cublas @ NVIDIA GeForce GTX 1080 Ti b=40 {}
ompnn/clang-18/tiled @ NVIDIA GeForce GTX 1080 Ti b=40 {"blas": "tiled"}
ompnn/clang-22/openblas @ Intel b=40 {}
ompnn/gcc-14/cublas @ NVIDIA GeForce GTX 1080 Ti b=40 {}
ompnn/gcc-14/tiled @ NVIDIA GeForce GTX 1080 Ti b=40 {"blas": "tiled"}
syclnn/clang-22/auto @ NVIDIA GeForce GTX 1080 Ti b=40 {}
syclnn/clang-22/tiled @ NVIDIA GeForce GTX 1080 Ti b=40 {"blas": "tiled"}