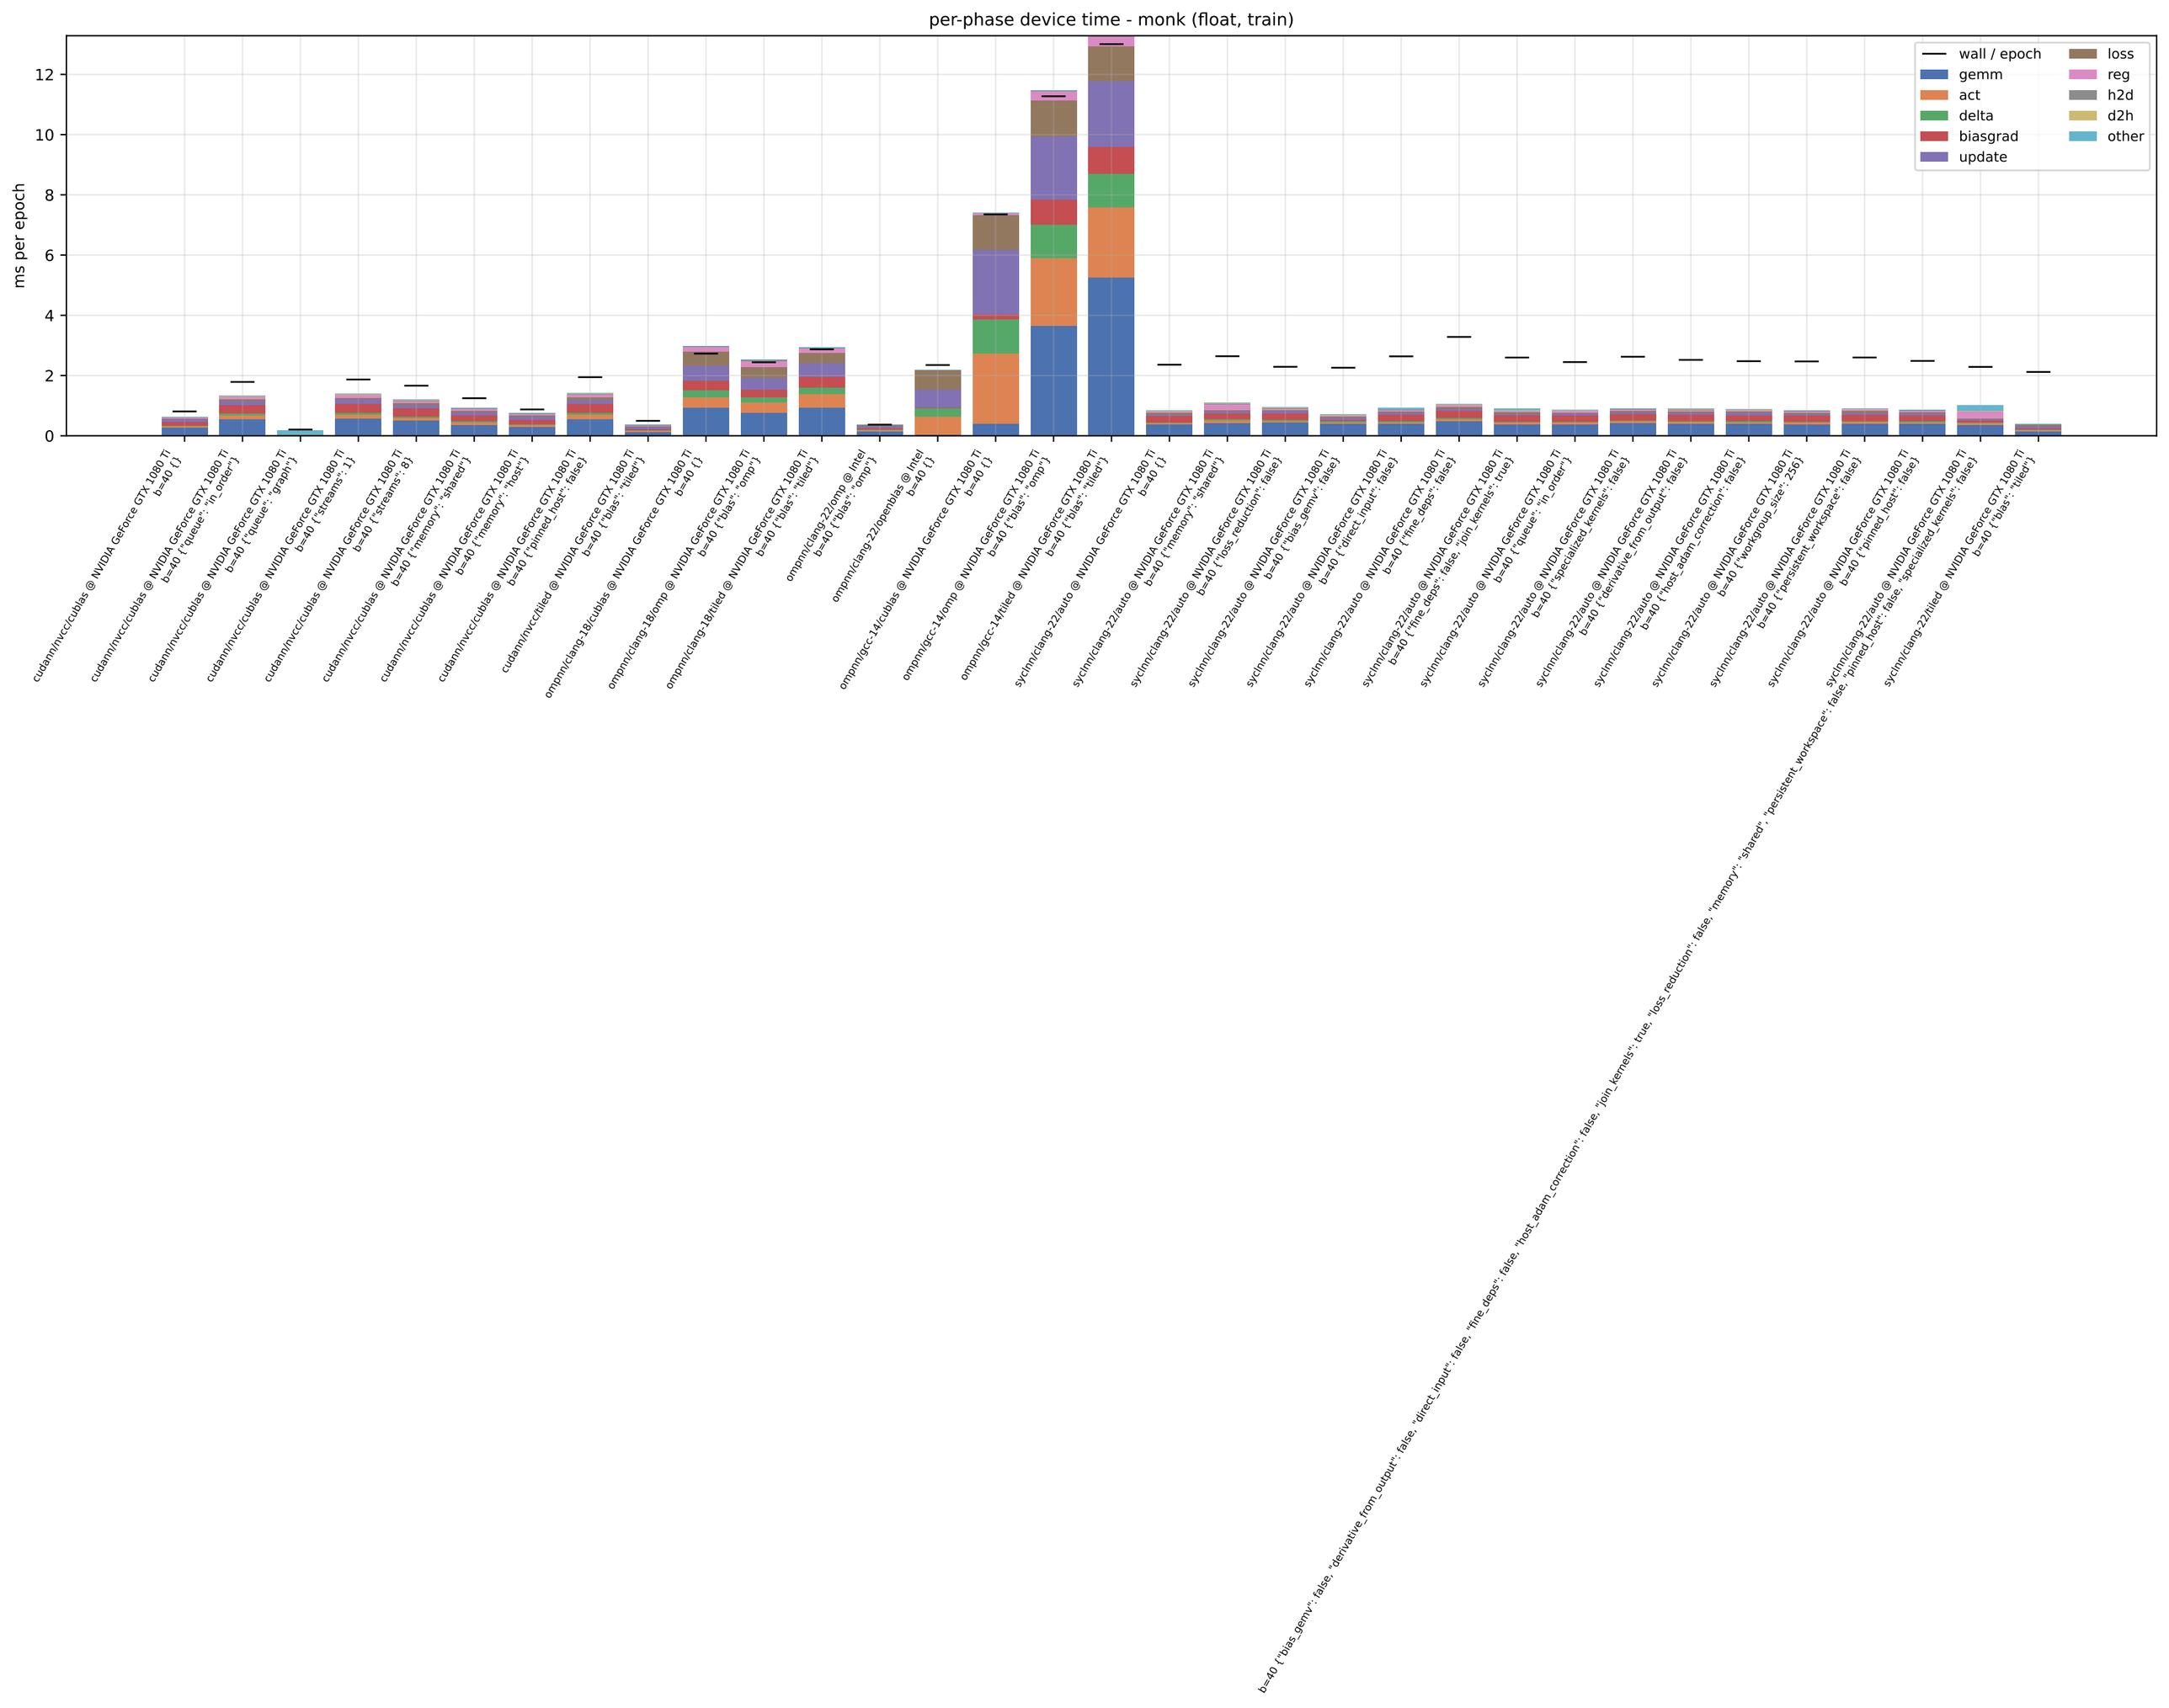

per-phase device time - monk (float, train)
ms per epoch
0
2
4
6
8
10
12
wall / epoch
gemm
act
delta
biasgrad
update
loss
reg
h2d
d2h
other
cudann/nvcc/cublas @ NVIDIA GeForce GTX 1080 Ti b=40 {}
cudann/nvcc/cublas @ NVIDIA GeForce GTX 1080 Ti b=40 {"queue": "in_order"}
cudann/nvcc/cublas @ NVIDIA GeForce GTX 1080 Ti b=40 {"queue": "graph"}
cudann/nvcc/cublas @ NVIDIA GeForce GTX 1080 Ti b=40 {"streams": 1}
cudann/nvcc/cublas @ NVIDIA GeForce GTX 1080 Ti b=40 {"streams": 8}
cudann/nvcc/cublas @ NVIDIA GeForce GTX 1080 Ti b=40 {"memory": "shared"}
cudann/nvcc/cublas @ NVIDIA GeForce GTX 1080 Ti b=40 {"memory": "host"}
cudann/nvcc/cublas @ NVIDIA GeForce GTX 1080 Ti b=40 {"pinned_host": false}
cudann/nvcc/tiled @ NVIDIA GeForce GTX 1080 Ti b=40 {"blas": "tiled"}
ompnn/clang-18/cublas @ NVIDIA GeForce GTX 1080 Ti b=40 {}
ompnn/clang-18/omp @ NVIDIA GeForce GTX 1080 Ti b=40 {"blas": "omp"}
ompnn/clang-18/tiled @ NVIDIA GeForce GTX 1080 Ti b=40 {"blas": "tiled"}
ompnn/clang-22/omp @ Intel b=40 {"blas": "omp"}
ompnn/clang-22/openblas @ Intel b=40 {}
ompnn/gcc-14/cublas @ NVIDIA GeForce GTX 1080 Ti b=40 {}
ompnn/gcc-14/omp @ NVIDIA GeForce GTX 1080 Ti b=40 {"blas": "omp"}
ompnn/gcc-14/tiled @ NVIDIA GeForce GTX 1080 Ti b=40 {"blas": "tiled"}
sycldnn/clang-22/auto @ NVIDIA GeForce GTX 1080 Ti b=40 {}
sycldnn/clang-22/auto @ NVIDIA GeForce GTX 1080 Ti b=40 {"memory": "shared"}
sycldnn/clang-22/auto @ NVIDIA GeForce GTX 1080 Ti b=40 {"loss_reduction": false}
sycldnn/clang-22/auto @ NVIDIA GeForce GTX 1080 Ti b=40 {"bias_genv": false}
sycldnn/clang-22/auto @ NVIDIA GeForce GTX 1080 Ti b=40 {"direct_input": false}
sycldnn/clang-22/auto @ NVIDIA GeForce GTX 1080 Ti b=40 {"fine_deps": false}
sycldnn/clang-22/auto @ NVIDIA GeForce GTX 1080 Ti b=40 {"fine_deps": false, "join_kernels": true}
sycldnn/clang-22/auto @ NVIDIA GeForce GTX 1080 Ti b=40 {"queue": "in_order"}
sycldnn/clang-22/auto @ NVIDIA GeForce GTX 1080 Ti b=40 {"specialized_kernels": false}
sycldnn/clang-22/auto @ NVIDIA GeForce GTX 1080 Ti b=40 {"derivative_from_output": false}
sycldnn/clang-22/auto @ NVIDIA GeForce GTX 1080 Ti b=40 {"host_adam_correction": false}
sycldnn/clang-22/auto @ NVIDIA GeForce GTX 1080 Ti b=40 {"workgroup_size": 256}
sycldnn/clang-22/auto @ NVIDIA GeForce GTX 1080 Ti b=40 {"persistent_workspace": false}
sycldnn/clang-22/auto @ NVIDIA GeForce GTX 1080 Ti b=40 {"pinned_host": false}
sycldnn/clang-22/auto @ NVIDIA GeForce GTX 1080 Ti b=40 {"pinned_host": false, "specialized_kernels": false}
sycldnn/clang-22/tiled @ NVIDIA GeForce GTX 1080 Ti b=40 {"blas": "tiled"}
sycldnn/clang-22/auto @ NVIDIA GeForce GTX 1080 Ti b=40 {"bias_genv": false, "derivative_from_output": false, "direct_input": false, "fine_deps": false, "host_adam_correction": false, "join_kernels": true, "loss_reduction": false, "memory": "shared", "persistent_workspace": false, "pinned_host": false, "specialized_kernels": false}