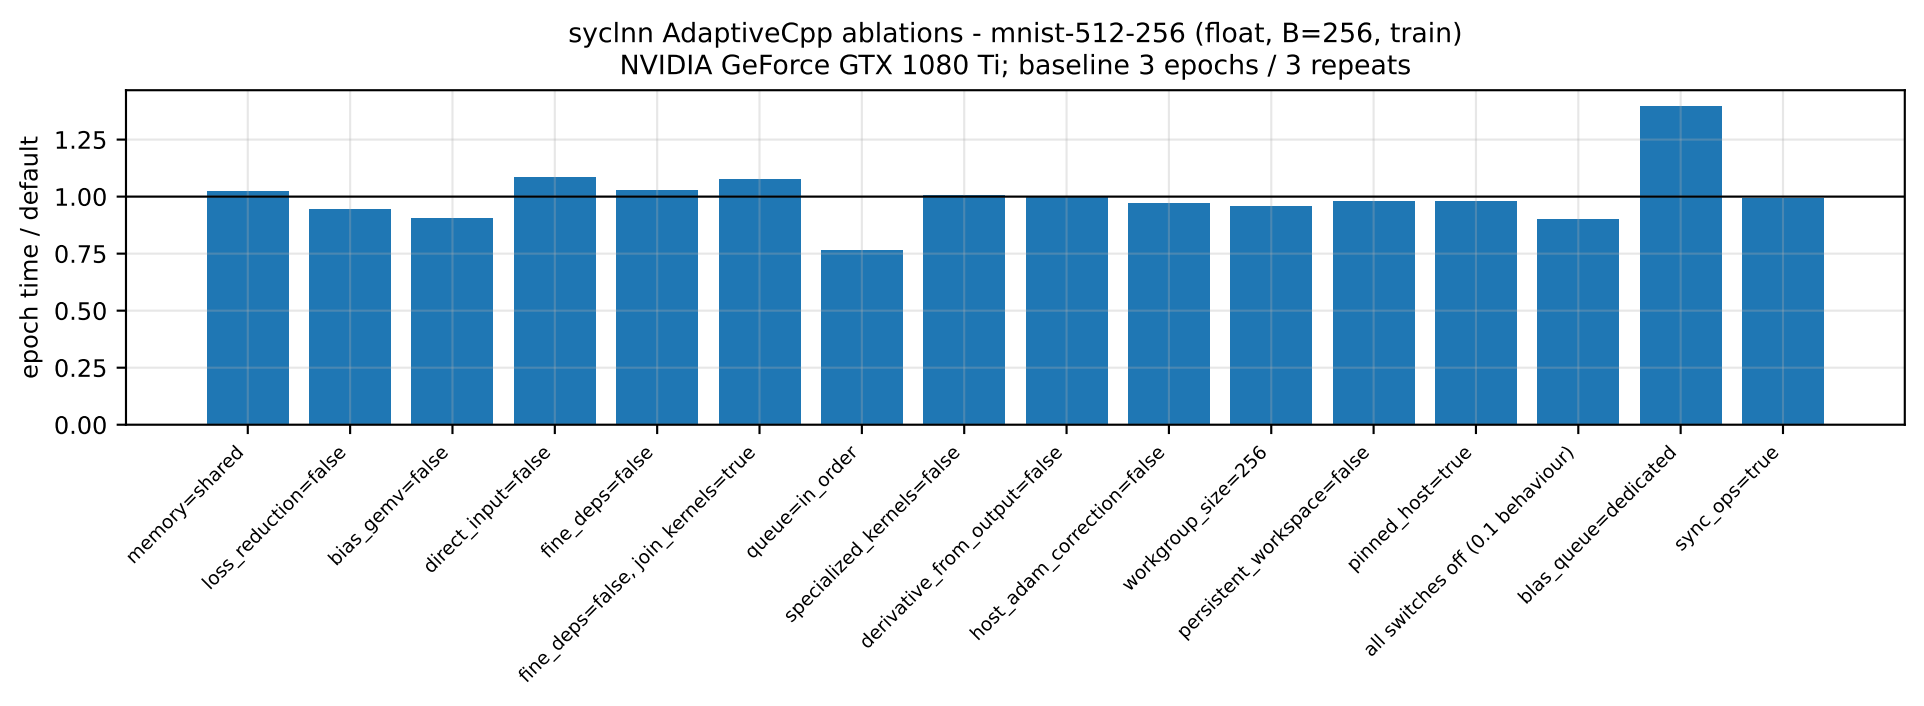
syclnn AdaptiveCpp ablations - mnist-512-256 (float, B=256, train)
NVIDIA GeForce GTX 1080 Ti; baseline 3 epochs / 3 repeats
epoch time / default
0.00
0.25
0.50
0.75
1.00
1.25
memory=shared
loss_reduction=false
bias_gemv=false
direct_input=false
fine_deps=false
fine_deps=false, join_kernels=true
queue=in_order
specialized_kernels=false
derivative_from_output=false
host_adam_correction=false
workgroup_size=256
persistent_workspace=false
pinned_host=true
all switches off (0.1 behaviour)
blas_queue=dedicated
sync_ops=true