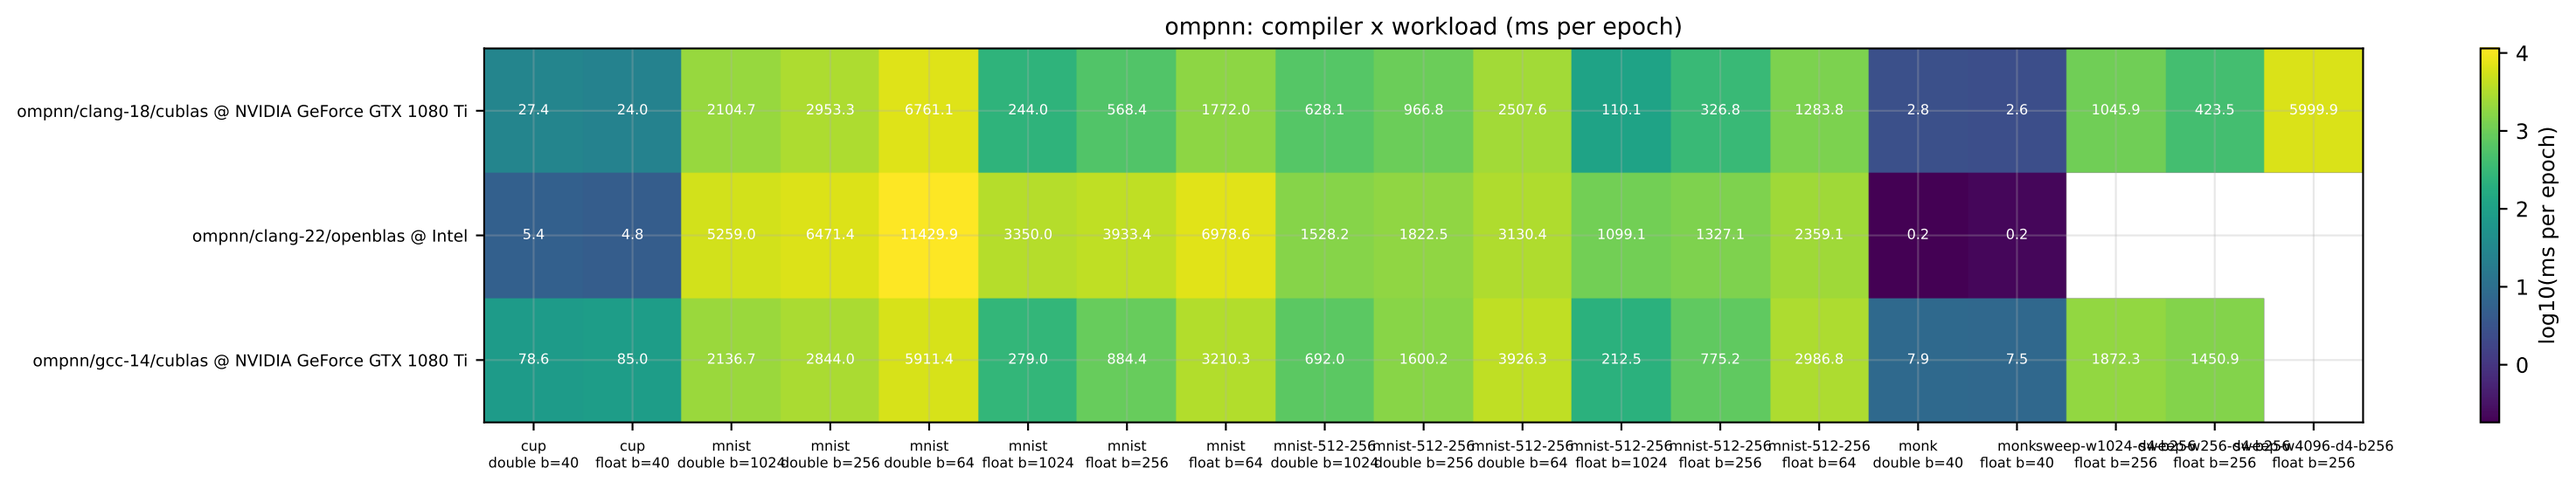
ompnn: compiler x workload (ms per epoch)
ompnn/clang-18/cublas @ NVIDIA GeForce GTX 1080 Ti
ompnn/clang-22/openblas @ Intel
ompnn/gcc-14/cublas @ NVIDIA GeForce GTX 1080 Ti
27.4
24.0
2104.7
2953.3
6761.1
244.0
568.4
1772.0
628.1
966.8
2507.6
110.1
326.8
1283.8
2.8
2.6
1045.9
423.5
5999.9
5.4
4.8
5259.0
6471.4
11429.9
3350.0
3933.4
6978.6
1528.2
1822.5
3130.4
1099.1
1327.1
2359.1
0.2
0.2
78.6
85.0
2136.7
2844.0
5911.4
279.0
884.4
3210.3
692.0
1600.2
3926.3
212.5
775.2
2986.8
7.9
7.5
1872.3
1450.9
cup double b=40
cup float b=40
mnist double b=1024
mnist double b=256
mnist double b=64
mnist float b=1024
mnist float b=256
mnist float b=64
mnist-512-256 double b=1024
mnist-512-256 double b=256
mnist-512-256 double b=64
mnist-512-256 float b=1024
mnist-512-256 float b=256
mnist-512-256 float b=64
monk double b=40
monksweep-w1024-d4-b256
monksweep-w256-d4-b256
monksweep-w4096-d4-b256
monk float b=40
monksweep-w1024-d4-b256 float b=256
monksweep-w256-d4-b256 float b=256
monksweep-w4096-d4-b256 float b=256
log10(ms per epoch)
4
3
2
1
0