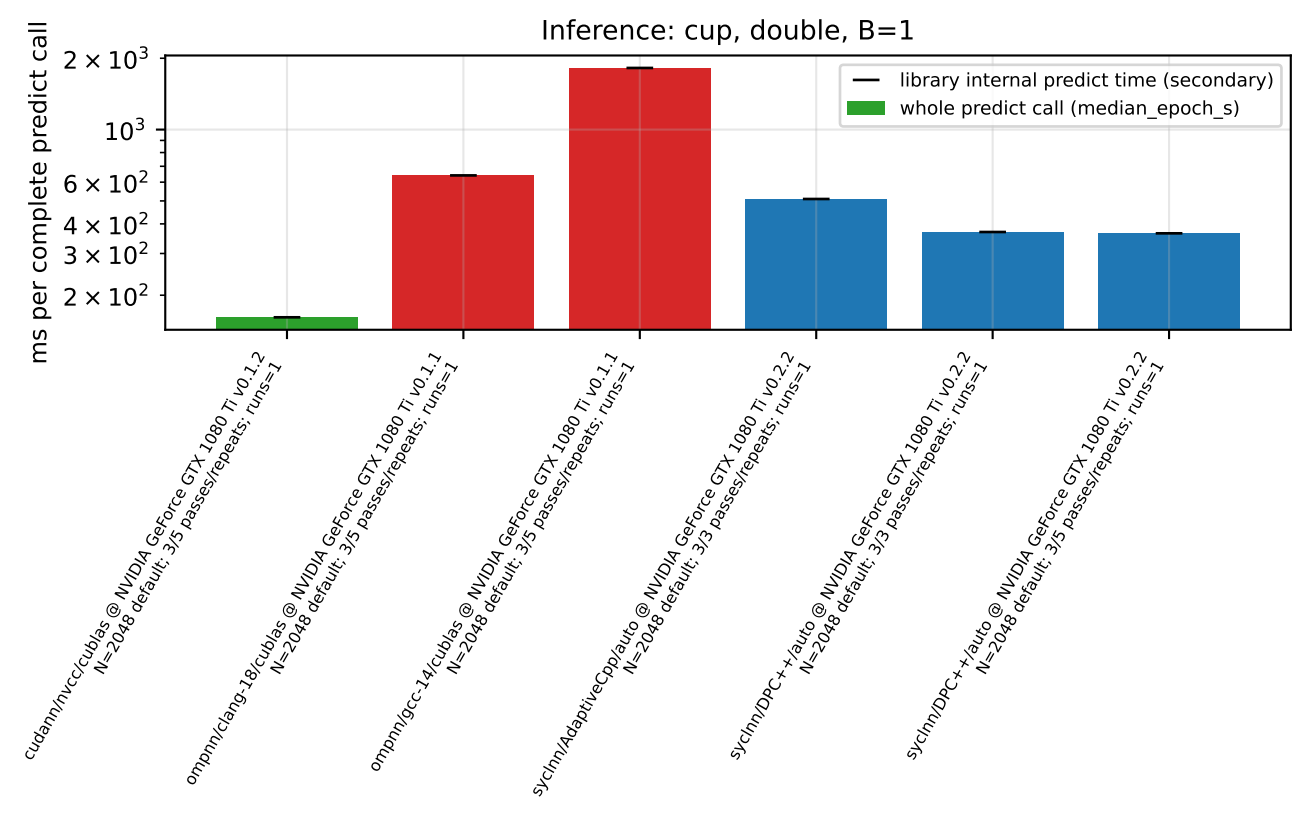
Inference: cup, double, B=1
ms per complete predict call
2 × 10^3
10^3
6 × 10^2
4 × 10^2
3 × 10^2
2 × 10^2
library internal predict time (secondary)
whole predict call (median_epoch_s)
cudann/nvcc/cublas @ NVIDIA GeForce GTX 1080 Ti v0.1.2 N=2048 default; 3/5 passes/repeats; runs=1
ompnn/clang-18/cublas @ NVIDIA GeForce GTX 1080 Ti v0.1.1 N=2048 default; 3/5 passes/repeats; runs=1
ompnn/gcc-14/cublas @ NVIDIA GeForce GTX 1080 Ti v0.1.1 N=2048 default; 3/5 passes/repeats; runs=1
syclnn/AdaptiveCpp/auto @ NVIDIA GeForce GTX 1080 Ti v0.2.2 N=2048 default; 3/3 passes/repeats; runs=1
syclnn/DPC++/auto @ NVIDIA GeForce GTX 1080 Ti v0.2.2 N=2048 default; 3/3 passes/repeats; runs=1
syclnn/DPC++/auto @ NVIDIA GeForce GTX 1080 Ti v0.2.2 N=2048 default; 3/5 passes/repeats; runs=1