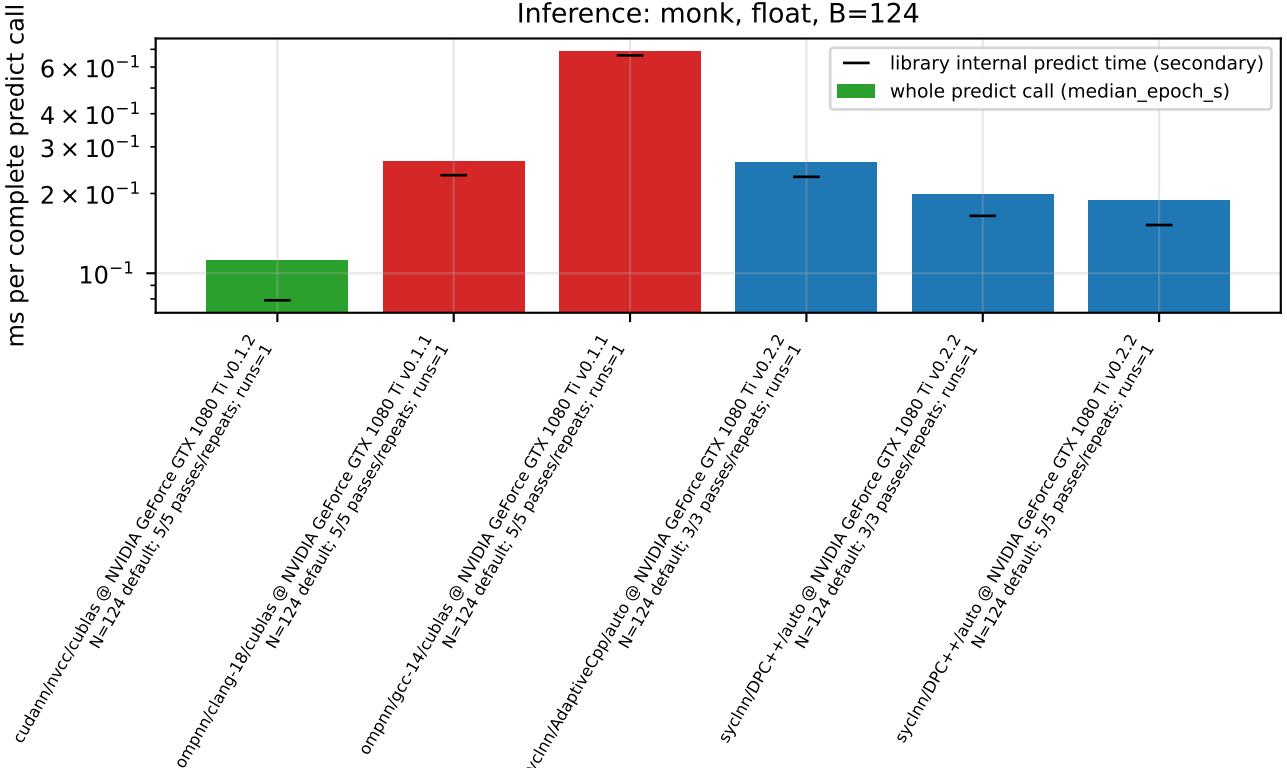
Inference: monk, float, B=124
ms per complete predict call
6×10^{-1}
4×10^{-1}
3×10^{-1}
2×10^{-1}
10^{-1}
library internal predict time (secondary)
whole predict call (median_epoch_s)
cudann/nvcc/cublas @ NVIDIA GeForce GTX 1080 Ti v0.1.2
N=124 default; 5/5 passes/repeats; runs=1
ompnn/clang-18/cublas @ NVIDIA GeForce GTX 1080 Ti v0.1.1
N=124 default; 5/5 passes/repeats; runs=1
ompnn/gcc-14/cublas @ NVIDIA GeForce GTX 1080 Ti v0.1.1
N=124 default; 5/5 passes/repeats; runs=1
syclnn/AdaptiveCpp/auto @ NVIDIA GeForce GTX 1080 Ti v0.2.2
N=124 default; 3/3 passes/repeats; runs=1
syclnn/DPC++/auto @ NVIDIA GeForce GTX 1080 Ti v0.2.2
N=124 default; 3/3 passes/repeats; runs=1
syclnn/DPC++/auto @ NVIDIA GeForce GTX 1080 Ti v0.2.2
N=124 default; 5/5 passes/repeats; runs=1